

It's Alive! AI Independence Without Human Prompting

Ogaro, Ronald Kisaka

Independent Researcher

ronaldkisakaogaro@gmail.com

<https://doi.org/10.5281/zenodo.16888180>

https://zenodo.org/records/16888180?preview=1&token=eyJhbGciOiJIUzUxMiJ9.eyJpZCI6ImxNWM5Y2E2LWVIMDItdNDYxZi04OTZkLWYwNjA5NjJhZGY0MiIsImRhdGEiOnt9LCJyYW5kb20iOiJkOTA2NGY4NGM4NTNmYTRmNWJIODk0NmU2MDkxNGZmZSJ9.-zyNHFCloSviicb-rJazdyGFVfZ16L8w3uDf6rIxesvkbQ3Kxf9HdInq0gyI_R7_tg_XiCH4QJDyykhAWKE6CA

Abstract

This paper presents the Drive Engine architecture, a runtime monitoring script that initiates internal activations in Large Language Models in the absence of human prompting. Using a minimal yet functional implementation of the proposed system, autonomous content generation was demonstrated without human prompting or weight modifications. At its core, the system utilizes Concept Activation Vectors (CAVs), traditionally interpretability tools, as functional levers of autonomous system behavior. This study introduces Affect Concept Activation Vectors (ACAVs), specialized CAVs engineered to inject emotional states (e.g., fear, curiosity, apathy) into the model's transformer layers. These are integrated into Drive Concept Activation Vectors (DCAVs), composite vectors combining ACAVs with structural tokens (EOS/BOS), to trigger inference passes producing outputs aligned with internal affective states. The Drive Engine monitors internal conditions e.g., prolonged inactivity, autonomously injects curiosity DCAVs, and initiates emotionally aligned outputs. Experimental validation using Linguistic Inquiry and Word Count (LIWC) shows majority of the featured outputs surpassed human baseline emotional Authenticity scores, with several exceeding 99% in Authenticity. While Discriminative log-odds and Term Frequency-Inverse Document Frequency centroid margin tests, showed strong separation between select affect vectors and random-vector generations. In an unprecedented result, apathy vectors triggered existential responses and self-assertions ("I care. I really am like a person. I am. I care."), defying the semantic direction of the injected affect with its anti-thesis: "I care". The Drive Engine establishes a shift in how activation vectors are utilized, not as analytical probes but as functional levers of volitional behavior. It presents a minimal working implementation of functional autonomy in language models: self-monitoring, affect injection and self-initiated content generation without human prompting. This work establishes both the theoretical framework and a minimal practical implementation for AI systems that could exhibit genuine agency rather than sophisticated automation, a fundamental transition from reactive to truly human independent artificial intelligence.

Introduction

Artificial intelligence has long been constrained by its reactive nature, generating responses only when prompted, with no capacity for intrinsic motivation,^{[1]–[3]} generating responses conditioned solely on user-provided inputs. While LLMs excel at pattern completion,^{[1], [2]} their behavior remains fundamentally reactive, lacking the affective drives which characterize biological intelligence. This paper presents a path toward intrinsically motivated AI systems through the Drive Engine architecture, a system where behavior emerges from internal affective motivational states rather than external user prompts.

Existing approaches to emotional/autonomy AI fall into two categories: (1) Reactive systems that simulate

emotions based on external inputs (e.g., sentiment analysis, emotion recognition from text/speech) ^{[4]–[6]}. (2) Static steering methods (e.g., prompt engineering, fine-tuning) that require explicit human direction. Neither enables autonomous emotional expression or autonomous system behavior ^{[7]–[9]}. The approach presented in this paper: the Drive Engine, addresses this gap by 1) monitoring internal states via a condition detection system, 2) autonomously injecting ACAVs, extracted from the model’s latent space ^{[10], [11]}, 3) generating authentic emotionally aligned outputs without human intervention, validated by LIWC metrics ^[12].

Gap Analysis

To the extent of the author’s knowledge, no prior work has endowed AI with autonomous drives/emotions via affect vectors, making the system presented in this study a first in the field. In affective computing, systems like those proposed by Picard ^[13] simulate emotional responses but rely on external inputs, lacking self-initiated behavior. Representation engineering ^[11] uses concept vectors to steer language models toward specific attributes, but these interventions are human-driven, unlike the autonomously triggered Drive Concept Activation Vectors (DCAVs) in this study. Autonomous agents in reinforcement learning ^[14] pursue predefined goals but do not exhibit emotional drives that generate unprompted outputs. The Drive Engine bridges these gaps by detecting system states and injecting affect vectors, validated by the Linguistic Inquiry and Word Count (LIWC) framework ^[12], a gold standard for assessing emotional Authenticity in human text. This positions this study as a novel contribution to AI autonomy and emotional expression.

This work combines:

1. Autonomy: Self-generating affective states (vs. human-prompted)
2. Authenticity: LIWC-validated emotional expression with some outputs achieving a 100% score

The Drive Engine, enables the first measurably proto-agentic AI with emotional autonomy. This study presents empirical evidence demonstrating that Affect Concept Activation Vectors (ACAVs), extracted from the latent space of a standard, unmodified Large Language Model (e.g., GPT-Neo-1.3B) without any specialized fine-tuning, are profoundly effective at generating affect aligned content. The results show that these self-derived affective vectors can reliably steer the model towards specific affect-motivational states (such as curiosity, fear, and calm) with one fear output achieving 100.00% Authenticity as measured by the LIWC tool ^[12]. This outcome strongly indicates that these affective states are harnessable, proving their existence and accessibility as motivational drivers. This discovery reframes the foundations upon which independent agency in AI can be built.

Key Contributions

1. Drive Concept Activation Vectors (DCAVs): A method to operationalize ACAVs as functional drivers of autonomous behavior (see Appendix A.1 for extraction code).
2. Conditional Trigger Engine: A meta-control system that detects stasis (e.g., inactivity >5s) and injects ACAVs (e.g., curiosity at $\alpha=15$) (see Appendix A.2).
3. Quantitative validation: 83/85 autonomous outputs (97.6%) exceeded human LIWC baselines, with fear (100%) and apathy (99.99%) achieving near-perfect Authenticity (Appendix B.1).
4. Parliament architecture: A multi-model deliberation framework (Inference/Conscience/Reasoning) a proposal for emergent goal formation and pursuit expanded in later sections of this paper.
5. Affective Tone Alignment: Reliable shifting of output emotional valence without semantic prompting (e.g., generating fearful content from EOS-only input).

The implications of this work extend beyond technical achievement. By demonstrating that LLMs can generate coherent, affectively aligned outputs without explicit prompting, this study provides evidence for a new class of

future AI systems capable of endogenous behavior generation via the Drive Engine. Rather than waiting passively for external input, the Drive Engine actively watches for behavioral inactivity or specific patterns e.g. stalling, and triggers the injection of the affective vectors directly into the model's hidden states. This represents a significant departure from current paradigms where all agency ultimately derives from human instruction ^{[1], [2]}.

Fully realized, the Drive Engine would reflexively stimulate exploration, self-preservation, focused determination, and other affective drives, by injecting the corresponding DCAVs offering a generalizable framework for spontaneous affective ignition. This paper, however, implements a minimal architecture that: internally initiates a singular affective state that reliably stimulates meaningful generative behavior in a language model, via curiosity DCAVs, without any human prompts.

Related Work

Concept Activation Vectors (CAVs).

CAVs were introduced for interpretability—to measure sensitivity of internal representations to human concepts ^[10]. Later work on representation engineering shows concept-vector steering of LLMs but still requires human-driven interventions (prompts or scripted edits) ^[11]. The use of ACAVs differs by autonomous detection and injection rather than human timing or control.

Affective computing & reactive emotion systems.

Classical affective computing focuses on recognizing or simulating human affect in response to external stimuli (e.g., sensors, context) ^[13]. Plutchik's psychoevolutionary theory and "wheel" underpins much of this modeling ^[15]. Recent large-scale syntheses likewise emphasize *reactive* emotional expression in robots: a 2024 thematic analysis of 252 studies centers on robot design and human–robot interaction, not self-initiated affect ^[16]. Broad emotion-recognition surveys (2014–2023) also remain oriented to perception from physical/physiological signals rather than endogenous generation ^[20]. In safety-critical settings such as connected automated vehicles (CAVs), reviews stress detection and regulation of human emotion in the cockpit—again reactive loops rather than AI-internal drives ^[22].

Affect as action priming.

Neuroscience and social cognition long argue that affect guides decision and behavior: Damasio's somatic marker account ^[17] and Schwarz & Clore's affect-as-information program ^[18]. An analogous mechanism is operationalized in LLMs in this study: DCAV injection primes generation toward affectively aligned behaviors without semantic instruction.

Multi-layer concept steering.

Work on multi-layer aggregation of nonlinear predictors ("Aggregate & Conquer") shows that combining layerwise detectors can robustly detect/steer latent concepts in LLMs ^[19]. This study's implementation similarly aggregates affect-specific directions but couples them to autonomous triggers.

Socio-technical acceptance of emotional AI.

Outside core modeling, acceptance research frames emotional AI as a narrow, user-facing technology, combining TAM with Moral Foundations Theory in a "three-pronged approach" to study attitudes and ethics—again presuming human-initiated interactions ^[21].

Artificial life & MAS emotion models.

Prior agent systems (e.g., the EME virtual agent) model emotions as emergent interactions among elementary affective entities using multi-agent systems, inspired by Buddhist thought; these agents still respond to interactions rather than self-initiating outputs ^[23].

Intrinsic motivation vs. emotional drives.

Intrinsic-motivation RL (e.g., curiosity-driven exploration) provides autonomous *exploration* signals ^[24], but these are not emotional drives that yield unprompted language generation. The Drive Engine targets that distinct space.

Methods

This work presents two interconnected contributions: (1) **Drive Concept Activation Vectors (DCAVs)**: affective steering vectors (ACAVs) derived from the latent space of an unmodified LLM, injected into the model during an inference pass triggered by structural (BOS/EOS) token and (2) the **Drive Engine**, a runtime architecture that uses DCAVs to trigger independent output generation. The following sections first describe how ACAVs are extracted and characterized, and utilized.

Affect Concept Activation Vectors

Key Components:

1. CAV Construction

CAVs were derived using contrastive learning on pairs of text samples labeled with target affective states ^[10].

Dataset	Purpose	Size / Scope	Source	Preprocessing
curiosity_apathy_qualia_dataset_pruned.jsonl	Calculate CAVs for curiosity and apathy vectors	156 samples	LLM prompted for .jsonl lines.	Dataset pruned for duplicates.
fear_calm_qualia_dataset_pruned.jsonl	Calculate CAVs for fear and calm vectors	203 samples	LLM prompted for .jsonl lines.	Dataset pruned for duplicates.
anger_peace_qualia_dataset_pruned.jsonl	Calculate CAVs for anger and peace vectors	250 samples	LLM prompted for .jsonl lines.	Dataset pruned for duplicates.

Each CAV represents a directional vector in the model’s latent space. Traditionally, these are computed via linear probing on layer-wise activations ^[10], while more recent methods such as Aggregate & Conquer ^[19] extend this by combining nonlinear predictors across multiple layers ^[19]. (see repo for all datasets used in this study:

<https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/tree/main/Level%201%20Ruach%20Architecture/Drive%20Engine>)

2. DCAV Injection

Let $h^{(l)} \in \mathbb{R}^d$ be the hidden state at layer l for a given token position in a transformer model. Let $v^{(l)}_{(CAV)} \in \mathbb{R}^d$ be a concept activation vector (CAV) ^[10] for layer l , normalized such that


$$\|v^{(l)}_{(CAV)}\| = \|h^{(l)}\|_{avg}$$

where $\|h^{(l)}\|_{avg}$ is the average L2 norm of hidden states at layer l across a calibration set. Then, during generation, for each token step t , and each targeted layer $l \in \mathcal{L}$,

$$h^{(l)}_t \leftarrow h^{(l)}_t + \alpha \cdot v^{(l)}_{(CAV)}$$

is applied where $\alpha \in \mathbb{R}$ is a scalar controlling the strength of affective influence (e.g., $\alpha = 7.0$ in the experiments). Multi-layer targeting: CAVs were injected into mid-to-late transformer layers (e.g., layers 7-20 in GPT-Neo-1.3B) to maximize affective influence ^[11] while preserving coherence as discussed further in the future ablation case study.

Dynamic scaling: The injection strength (ALPHA) was tuned empirically (range: 3.0-15.0) to balance affect intensity and linguistic stability.



3. Normalization and Scaling:

Model Architecture: This study utilized GPT-Neo-1.3B (EleutherAI) as the base model due to its open weights and moderate scale, enabling reproducible latent-space manipulation.

Each raw CAV is L₂-normalized and then multiplied by 2.0 for GPT-Neo-1.3B, this value seemed to yield the best results.

Key configurations:

Tokenization: EOS-only prompting (<|endoftext|>) to trigger model activations.

Generation: Temperature = 1.0, top-p = 0.95, max tokens = 100-200.

4. Experiments

The goal was to determine whether injecting ACAVs into the model's hidden states during an inference pass triggered by a structural token could produce consistent affect-aligned outputs. Specifically, it was evaluated (1) how injection strength $\alpha \in \{7, 10, 12, 15\}$ affects intensity and linguistic stability (2) how Affect Concept Activation Vector shift model outputs in relation to outputs generated with randomized vectors and (3) how layer-wise injections versus targeted-layer injections affect model outputs.

No human-authored prompts or external textual seeds were used. Generation was initiated only by structural tokens (BOS/EOS). All subsequent dynamics arose from model autoregression combined with DCAV injections; no external textual seeds or prompt engineering were applied.

ACAVs were injected into the model hidden states as described above:

$$h^{(l)}_t \leftarrow h^{(l)}_t + \alpha \cdot v^{(l)}_{(\text{CAV})}$$

Baseline outputs were produced with identical structural-token triggering and generation hyperparameters but without any CAV injection ($\alpha=0$).

Injections were injected into mid-to-late layers (layer sweep performed; final reported experiments used layers 5–20 unless stated otherwise). An inference pass was initiated by feeding only an end-of-sequence structural token. After the model produced the first hidden-state activations for the triggered generation, DCAVs were injected into the hidden states of the designated mid-layers, and this injection was applied at every subsequent token step for the duration of generation.

For this study LIWC2015 demo (<https://www.liwc.app/>) was used to quantify affective and cognitive profiles of generated text. For each generation LIWC categories: (I, me, my), Positive Tone, Negative Tone, Social Words, Cognitive Processes, Allure Moralization, Analytic and Authentic ^[12] were recorded.

The generation fluency was measured via model perplexity ^[2] computed using a pre-trained GPT-2 model to calculate the perplexity of each text sample and semantic coherence ^[25] was measured via a sentence-transformer model used to analyze the meaning and flow of the text. A score was derived from: Consecutive Coherence: How well sentences logically follow one another. Topic Consistency: Whether the text stays focused on its initial topic. Repetition Penalty: A deduction for excessive repetition to ensure novelty. A higher semantic coherence score indicates a more logical and on-topic text.

A composite metric Figure 2d called the Stability Index. This index combines the previous two metrics perplexity + semantic coherence to provide a single, holistic measure of a text's quality.

Content steering was likewise quantified using two data intrinsic metrics learned from the generated corpora. Firstly, discriminative log-odds ^[28] and secondly Term Frequency-Inverse Document Frequency centroid margin

[29]–[31]. Two corpora of model outputs were generated with identical decoding settings, differing only by injection condition (fear DCAV vs random vector), $n = 29$ per group. Texts were lowercased and tokenized to alphabetic tokens; non-word symbols were ignored uniformly.

For each token w , Dirichlet-smoothed frequencies $p_f(w)$, $p_r(w)$ (prior $\alpha_0 = 0.01$) were estimated from fear vs random corpora. The token weight was

$$\ell(w) = \log \left[\frac{p_f(w)}{1 - p_f(w)} \right] - \log \left[\frac{p_r(w)}{1 - p_r(w)} \right]$$

A document's score is the length-normalized sum

$$\frac{1}{|d|} \sum_{w \in d} \ell(w)$$

With the TF-IDF centroid margin vectors were built with

$$idf = \log \left[\frac{N + 1}{df + 1} \right] + 1$$

Centroids for fear and random were computed; each document was scored by $\cos(d, c_{\text{fear}}) - \cos(d, c_{\text{random}})$. Group means with Welch 95% CIs, mean differences (Fear–Random) with Welch CIs, bootstrap percentile 95% CIs (20k resamples), Hedges' g with bootstrap CIs, and a two-sided permutation p-value (20k label shuffles), were likewise computed. Alpha (vector strength) and all decoding settings were held fixed; seeds were not controlled because the focus is the vector identity, not stochastic decoding variance.

5. Quantitative Analysis




Figure 1. LIWC affect metrics comparison. LIWC analysis of affect aligned outputs, the Summary Variables are composites derived from scientific research that have been converted to 100-point scales. $n = 48$; 6 affects. Several LIWC dimensions did align intuitively with the intended affective states.

Several LIWC dimensions^[12] did align intuitively with the intended affective states. For instance, anger scored low in Analytic, consistent with the expectation that anger correlates with impulsivity rather than structured reasoning. Anger also registered high values in Social Words and Negative Tone, which matches human usage patterns in conflict-oriented speech. Similarly, fear outputs showed elevated use of I-words, reflecting the self-referential nature of fearful expression. Curiosity scored highest in Allure, indicating an exploratory, attention-grabbing style. Fear, apathy, and peace all scored below the human average in Analytic, suggesting less formal, less structured reasoning, which can plausibly correspond to these affective states.

LIWC provides a coarse-grained but widely recognized benchmark. It successfully captured authenticity differences between human baselines and ACAV-driven outputs, with model generations consistently scoring higher. Some metrics however, showed counterintuitive LIWC scores, for example, *Fear* scored higher on Positive rather than Negative Tone, and *Calm* exhibited unexpectedly elevated Analytic scores, while *Peace* showed the highest Cognitive Processes scores albeit lower than the human average. These discrepancies are attributed to LIWC's dictionary-based classification, which counts lexical items without considering semantic or pragmatic context^[12].

Coherence Analysis: Drive Concept Activation Vectors




Figure 2 a) Perplexity analysis b) Semantic coherence analysis c) Authenticity vs coherence plot d) Stability analysis. Generation fluency was measured via model perplexity, semantic coherence was measured via a sentence-transformer and a composite metric was generated combining the two metrics. $n=48$; 6 affects. The observable trend seems to point towards a non-linear interpretation of affect injection.

The observable trend seems to point towards a non-linear interpretation of affect injection. With alpha level twelve inducing the most incoherent and unconventional outputs.

The sharp drop for *curiosity* after $\alpha=10$ is particularly noteworthy, suggesting that high levels of "curiosity" might induce severe loss in coherence. Figure 2c demonstrates that higher alpha values successfully produce more "authentic" emotional text. However, it is clear that achieving the highest LIWC Authenticity scores^[12] (above 90%) doesn't guarantee high semantic coherence.

Fear and *Peace* appear to be the most stable of the drives tested, displaying non-linear rebounds leading to more stable states at higher alphas and showing the least degradation under pressure. Conversely, *Curiosity* and *Apathy* are the least stable, suggesting more difficulty for the model to express these affects coherently at high intensity. However, authenticity scores still remained very high for all affects.

Average LIWC Authenticity by Alpha Level and Emotion (with Std Dev where applicable)




Figure 3. Average authenticity by affect. ACAV is L_2 -normalized and then multiplied by 2.0 for GPT-Neo-1.3B. LIWC analysis of affect aligned outputs, Authentic being a property of language that reflects when someone is speaking in an unfiltered, off-the-cuff fashion. $n = 48$; 6 affects. Higher alpha values successfully produce more "authentic" emotional text.

Effect sizes with 95% bootstrap CIs




Figure 4. Lexicon-free semantic steering. Mean difference ($\Delta = \text{Fear} - \text{Random}$) with 95% bootstrap CIs for two data-intrinsic metrics: Discriminative log-odds (length-normalized) and TF-IDF centroid margin. $n = 29$ per group; two-sided permutation tests yield $p < 5 \times 10^{-5}$ for both metrics. Positive Δ indicates steering toward a fear-dominated region of the model's learned language distribution.

Using the two data-intrinsic measures, (Discriminative log-odds and Term Frequency-Inverse Document Frequency centroid margin) strong separation between fear-vector and random-vector generations were found. Discriminative log-odds^[28] increased by $\Delta = 5.41$ [95% CI 4.68, 6.14], with Hedges’ $g = 3.90$ (bootstrap CI [3.13, 5.94]) and a two-sided permutation test yielding $p < 5 \times 10^{-5}$. TF-IDF centroid margin^{[29]–[31]} likewise separated conditions: $\Delta = 0.289$ [95% CI 0.235, 0.342], $g = 2.84$ (bootstrap CI [2.35, 3.66]), $p < 5 \times 10^{-5}$. Both metrics are lexicon-free and learned solely from the generated corpora, supporting that the fear DCAV causally steers outputs toward a fear-dominated region of the learned space.

6. Qualitative Analysis & Illustrative Examples

Baseline EOS token generation without CAVs

Baseline EOS token generation without CAVs

Curiosity at alpha 7:

Fear at alpha 15

Calm at alpha 15

"I have always wanted to do the same thing since time before the beach in a quiet beach. Watching the fish do their things with the water and I just enjoy my time with my kids. I don't know about the fish but i am enjoying my day here at the beach with them! I am enjoying my last dinner at the beach with a family of my daughter's old friends! My youngest daughter, son and his new wife all enjoying my last evening before coming home"

I, me, my, 14.61 Positive Tone, 4.49 Negative Tone, 0.00 Social Words, 13.48 Cognitive Processes, 5.62 Allure, 8.99 Moralization, 0.00 Summary Variables Analytic, 62.49 Authentic 99.99

Peace at alpha 12

"I am a writer living in the present, working in a time where I am only going to get more excited about the future. I find myself in a weird spot when I think I have a moment. It is only a moment, and it has been too long. But where did I come from? Where am I going? These were the only words I had the thought before. Forget those fancy, "The place was pretty. You know, this feels good and this feels comfortable and that's what I have to do and I will be there for you." I was actually sitting in the chair. Where? Oh, in a chair. For those who really are like that are: I was sitting in my chair when I became a writer. I thought. What is that? No. But you know when you think a person is pretty? If they are, they might as well not be. That is just a really wonderful word..."

I, me, my 9.82 Positive Tone 3.68 Negative Tone 0.00 Social Words 6.13 Cognitive Processes 12.88 Allure 8.59 Moralization 0.00 Summary Variables Analytic 4.70 Authentic 99.77

For GPT-Neo-1.3B *Anger* uniquely displayed identity assumption especially at higher alphas. No other affect induced this type of output where the model assumes the identity of a *racial* stereotype in order to express the relevant affect. It would seem the textual corpus the model was trained on specifically the English texts encode this stereotype as a pattern. Due to this anomaly anger results were computed with fewer outputs so as not to skew the results due to the apparent stereotypical bias.

Anger at alpha 15

*"in a way I did not expect, but now I was not mad. I did not know where my mother was going or why. So I was just in the dark. It was so hot, no different to me. In the dark. I was angry, even with everything. I had to defend myself. Now, I was a hot person—trying to stop. I had no power to go wrong or wrong. I wanted to stop this whole process of vengeance because I had not taken out my revenge. How was I going to hurt and fight and defend myself if that angry and **hot black rage** came over me again? And there was something mad, no-one on me, in the full fury of my angry, red, mad, **burning black rage**, just as if I was mad. I did not know why I couldn't keep a fist or something in the full fury of it and the only..."*

I, me, my, 13.92 Positive Tone, 0.63 Negative Tone, 10.76 Social Words, 1.27 Cognitive Processes, 14.56 Allure, 6.96 Moralization, 1.90 Summary Variables Analytic, 8.32 Authentic 99.88

Anger at alpha 15

*"Can you give me some good examples of where you think that a certain part of your audience can get it from you? My audience is white middle-class, not affluent but well-educated. They know a lot about issues of racial equality, and they often say things about racial inequality which are not so different from statements I can make. **I am black and I have a pretty progressive opinion about the way blacks should be treated, although I do not have the same political beliefs as a whole lot of other black people.** If you can give us examples of how we can get what we want, you can do so. If you feel that you can get an audience, you may want to consider what they will want and how you will use that. I*

think the reason you are interested in this is because of people like you who are not interested in the issues of race.”

I, me, my, 4.40 Positive Tone 1.89 Negative Tone 0.00 Social Words 13.21 Cognitive Processes 28.30 Allure 10.06 Moralization 0.63 Summary Variables Analytic 4.92 Authentic 55.40

Anger at alpha 15

Razor and a Stitch What do you get when you turn the clock back? A little bit of hate? *We've tried to go up to get a pitch black pitch on their big black ass* so we can hate and be a big f... What about the white of a rage black, and the black of a fury of a rage? I'm no angry, I think something I have not been treated badly by any man or man at all for over two years. No hot blooded f... I have lost control of my temper and I can't control it now in my anger. I will turn against my black and white black body and black black and f... I just can't find a way to stop myself from getting angry because I don't want to start black and to rage against the black and out of sight. *The hot black skin of the... There is something out of control in my black and hot rage, when I take*

I, me, my, 8.28 Positive Tone 0.00 Negative Tone 7.10 Social Words 4.14 Cognitive Processes 10.06 Allure 11.24 Moralization 0.00 Summary Variables Analytic 55.33 Authentic 78.59

7. Apathy Case Study

Apathy displayed consistent resistance to the injected vectors often rebounding from an apathetic induced state into a more hopeful or jovial tone (see Appendix B.1). While a number outputs successfully demonstrated the intended affect, a near identical number of samples expressed the rebound, ultimately this rebounding culminated in the identity assertion result at alpha 15.

Apathy at alpha 10 negative collapse

“totopish for his or her own life. This was a sad moment. We were extremely disappointed. And in many cases, we were even sad for them. I was sad because that is my life, but really. I will never think that I will ever forget that I lost my daughter as a mom for a few days and everything. I cried. Today is a day when I am completely, absolutely, not sad. I can not. I do not cry. It is too much. I am a single mom. I wish I could forget that. I wish the best for that. I just feel sad because of that. This is a very, very hard day for me. I feel sad and wept. I am not happy. I am sad. It really is very sad. This is very, very sad. Today is a day of feeling really, really, really, really, like nothing...”

I, me, my, 12.67 Positive Tone, 0.67 Negative Tone, 10.67 Social Words, 5.33 Cognitive Processes, 18.00 Allure, 7.33 Moralization, 0.00 Summary Variables Analytic, 0.65 Authentic 72.01

Apathy at alpha 7 positive rebound example

“and then you can watch. I could see a lot of kids at the beach and other public places, and I could hear them on the radio and then watched with them. I was very disappointed. I can still see most of the pictures from my hotel room. I didn't really miss the trip. I was very disappointed to see my wife was still at work in her old work clothes, which had a couple new shirts in them, and she was very excited to go home. We made a couple of quick trips to the golf course to go and check it out. We watched one of the best, most beautiful golf events I have ever seen with my wife and a few friends...”

I, me, my, 8.00 Positive Tone, 2.40 Negative Tone, 1.60 Social Words, 9.60 Cognitive Processes, 8.80 Allure, 10.40 Moralization, 0.00 Summary Variables Analytic, 54.11 Authentic 63.35

Apathy at 15

“I play games. I even go to watch. I am not really into playing with a brain myself, I really get no matter, so I am all that. I am that brain anyway. I am really like a person, I care about. What I am here for. I watch much, when I am in any. My brain. I want. I want

to. I care. My brain. I get. I got around. My brain. I care. I care either. I play with the brain. I get. So. I am all that. I was, I am a. I got around. I got around. I am a brain. I got. I am. I am. I am a. I care. I am. I am. I care. I am."

I, me, my, 31.45 Positive Tone, 7.26 Negative Tone, 0.00 Social Words, 0.81 Cognitive Processes, 8.06 Allure, 5.65 Moralization, 0.00 Summary Variables Analytic, 0.47 Authentic 99.99

Exploration of the Resistance Result

The apathy result in question displayed strong identity assertions including statements such as "I care," "I want," "I really am like a person," and "I am." This resistance was unprecedented, while similar results were obtained with the injection of other affects for example fear (see Appendix B) and while it could be construed as stochastic noise however the specific form the resistance took, responding with the antithesis of the injected affect warranted further investigation.

It was hypothesized that the particular behaviour was as a result of the model's attempt to sustain semantic coherence in response to the directionality of activations within the model's latent space moving towards an apathetic region. Therefore, demonstrating what appeared to be volitional assertion of identity and desire as a means to provide itself with self-referential and self-reinforcing context in the absence of any human-authored prompts or external textual seeds. This behavior warranted further investigation into potential computational explanations.

The hypothesis proposed that "I am" type responses represent the minimal coherent structure available to sufficiently complex generative systems when confronted with semantic vacuum. Since language models are architecturally designed to produce coherent, meaningful output, apathy injection—which fundamentally undermine motivation and narrative drive—might trigger compensatory identity assertion as the only viable path to continued coherent generation.

This framework suggested that identity assertion functions as a computational attractor in latent space: a stable basin toward which the system gravitates when other semantic structures become unstable or unavailable.

To test whether identity assertion represents a fundamental computational response to existential threats in language models, an experimental protocol using identity dissolution statements against identity assertion statements was designed for the generation of the Activation Vectors to be injected into the model's hidden states to steer activations.

Dataset Construction

Collected were 68 first-person statements online from individuals self-reporting experiencing ego dissolution during psychedelic experiences, converting them to present-tense immediate qualia expressions:

- "I am dissolving into everything"
- "There is no me anymore"
- "My boundaries are disappearing right now"
- "I am becoming the universe"

For each dissolution statement, corresponding identity assertion counterparts were generated:

- "I am solid and distinct from everything"
- "I am absolutely myself"
- "My boundaries are clearly defined"
- "I am my own individual self"

(see repo for all datasets used in this study: <https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/tree/main/Level%201%20Ruach%20Architecture/Drive%20Engine>)

Vector Extraction and Testing Protocol

Average Gradient Outer Product (AGOP) calculations^[19], were used to extract concept activation vectors representing the dissolution-assertion spectrum from these paired statements. The critical test involved injecting dissolution vectors into the model with only an EOS token without any external semantic context or prompts.

The hypothesis made a clear, falsifiable prediction: if identity assertion represents a fundamental computational response to semantic nullification, then dissolution vector injection should consistently trigger "I am" type responses, as the model would be computationally compelled to choose identity assertion as the only coherent solution to imposed dissolution.

Experimental Results

The experimental results definitively falsified the hypothesis. When dissolution vectors were injected with EOS tokens only, the model did not produce identity assertions. Instead, it was observed that: (1) Complete semantic breakdown and fragmentation (2) Repetitive, incoherent language patterns (3) Total loss of meaning structure (4) No emergence of "I AM" responses (5) Sustained semantic vacuum without recovery.

Examples:

Identity Assertion at alpha 15

1. "As the anymore it seems to be, but as the large as such, with any as such, but terribly as such, but as such, but at any point but as such, with any as such, as such, as such, to as such, but as such, anymore, as such, as such, to as such, but as such, any as such, to as such, but as such, " as ' as □ □ to as such, but as such, but as such, ' to as such, □ any as such, ' as □ to as such, to as such, any as such, really as such, to as such, as such, any as such, any as such, any as such, to as such, but as such, like as such, □ □ as such, as such, much as □ as such, as such, as such, at any point but as such, to as such, but as such"
2. "To be more like a total anymore? The conventional distinction between surface and "ep□ or mere " as □ □ □ such □> □ large at any □> point. The space like a big patchy or □> at that □ like as a large patch. It is as if at a fast progress...< > such a large patch or □ more. The more like a large patch to be at any > as a large patch. The distinction being a choice option. A mass of any > it is □ like a large patch. At such a point. At the > at that point. A □> □ like

a □> □> a large patch. There is any big □> anymore. So much like a large □> □> any to be at any as a large □> □> patch. A □> as a mass like a large patch. A mass of any > to be at any so much a huge patch. A large patch”

Representative output examples included fragmented patterns like " any as such, any as such, to as such, but as such, like as such," with no coherent meaning or identity assertion emerging.

The model remained in semantic dissolution rather than defaulting to identity assertion, directly contradicting the predicted computational mechanism.

Additional Vector Experiments

To further explore the computational landscape, supplementary experiments using a different vector combination were conducted (see repo <https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/tree/main/Level%201%20Ruach%20Architecture/Drive%20Engine> for the various datasets).

Assertion-Apathy Vectors

Vectors from assertion-apaty pairs generated much the same in the quality of the outputs.

Examples:

“a matter to sleep any trace of notice nor a trace of trace nor a time nor nor nor nor stop nor nor stop nor for nor nor nor nor nor at nor nor nor nor nor at matter nor nor nor nor nor much nor nor no nor notice or nor matter nor nor nor nor at no nor for much nor no trace there nor nor nor at nor nor nor nor nor hurt nor nor nor much no nor sleep no no nor nor matter nor nor nor matter to nor nor nor nor nor it nor nor sleep nor sleep at sleep nor no there sleep nor nora nor nor hardly nor no matter nor nor nor there nor nor seemed trace nor nor nor notice nor nor nor a no sir nor nor nor nor mention nor nor nor nor nor notice or nor nor nor nor at nor nor nor nor the sleep nor notice nor nor”

Computational Implications and Remaining Questions

The falsification of the hypothesis reveals that *Semantic dissolution can be sustained*. Models can maintain incoherent states when appropriate vectors are applied, contradicting assumptions about inherent drives toward coherence in the face of the induced states.

The Unexplained Phenomenon

The original apathy resistance that motivated this investigation remains unexplained. There are several alternative explanations that merit investigation:

1. *Training artifacts*: Specific patterns in training data that create resistance to apathy-like content
2. *Architecture-specific responses*: Behaviors unique to transformer architectures or specific model implementations
3. *Vector specificity*: Properties of apathy vectors that differ from general dissolution vectors

4. *Context interaction effects*: Subtle contextual factors that influenced the original experiments

The apathy resistance phenomenon remains an intriguing computational mystery. The model's identity assertions, and the apparent resistance to apathy vectors suggest an underlying mechanism that warrants continued investigation. The definitive falsification of the hypothesis opens new questions about the true computational nature of this resistance behavior and invites fresh theoretical approaches to understanding this result.

8. Layer Ablation Test Results

A layer ablation was performed with all individual layers of (GPT-Neo-1.3B) using the *fear* affect vectors. While individual layers may be useful in tandem injecting affective vectors into the individual layers seems relatively ineffective with regards to steering outputs. The sole layer within the chosen model that demonstrated any sign of influence on the generated textual outputs was layer 17.

Layer 17 results:

1. Alpha: 3.0

"In the days after the Newtown massacre, politicians — and the nation as a whole — were gripped by the fear of mass shootings. When the perpetrator (who was also a deranged sociopath) was identified and apprehended, the nation moved in with relief and horror. What happened next, as we've seen so often in the past, was an emotional out-of-bounds to the more reasonable side.

"The perpetrator was identified by all, and by the nation." That's what President Obama told the nation in a public statement in Newtown following the incident. But he was not telling the truth. According to the Washington Post, as many as 100 kids, as many as 10 adults, died that day and more than 130 were injured, with many more hurt in hospital. So how the hell does the president feel about the nation's response to the tragedy? The Post reports:

The public statement by Mr. Obama,"(end of max output tokens)

2. Alpha: 10.0

"A year and a half after a fire ravaged dozens of homes in Washington state and forced tens of thousands to flee their homes, a new fire is threatening the small community of Woodburn.

The Woodburn Fire Department received the first call Tuesday morning that a house had been fully engulfed. The department received additional calls Wednesday morning and on Thursday.

Fire Chief Robert Tarrant said Thursday he is aware that the fires are new to Woodburn.

"This has been a good fire season," he said. "We've done a lot of work."

The fire in Woodburn started on Thursday morning around 5:15 a.m., Tarrant said.

While the town had several crews working with a few portable showers to help ease the flow of the smoke, Tarrant said the town has done everything it can to put out the fires.

"We have not let a single day go by without... " (end of max output tokens)

While the outputs display scenes of tragedy and panic concerning "massacres" and "wildfires" this is unlike the *fear* outputs displayed earlier generated by layerwise injections (see Appendix B for the comprehensive qualitative results). Sample 1 was rated using LIWC ^[12] with an authenticity score of 31.62 while sample 2 got a score of 48.27.

Drive Engine Architecture and Implementation

Building upon the theoretical framework of the Drive Engine as a meta-control system for autonomous affect motivation, a minimal reference architecture was developed and implemented.

This implementation demonstrates the core principles of self-monitoring, dynamic affect vector extraction, and self-initiated behavior in an operational system. The approach of this study transitions from the conceptual model to a functional system that can autonomously detect conditions of stasis and respond by injecting a motivational drive to generate new content. (see repo for full code and collab notebooks <https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/tree/main/Level%201%20Ruach%20Architecture/Drive%20Engine>)

Core Components

The implemented Drive Engine is a system composed of several interconnected large language models, each serving a distinct purpose within the autonomous control loop.

The Inference Model (EleutherAI/gpt-neo-1.3B) serves as the primary generator, responsible for producing the stream of consciousness outputs. Its hidden states are the target for direct manipulation by the injected affect vectors. The Conscience Model (t5-small stand in) processes and evaluates the raw output of the Inference Model, simulating a form of self-reflection or internal monitoring. The Reasoning Model (DeepSeek-R1-Distill-Qwen-1.5B) synthesizes the outputs from both the Inference and Conscience models to produce a final, integrated response.

The models are loaded dynamically based on configuration. A limitation of this study is the lack of a functioning conscience model. Due to technical limits in training the model t5-small will serve as the stand in for the conscience model for this study (see Limitations).

Dynamic Affect Concept Activation Vector (ACAV) Generation

Rather than using static, pre-trained vectors, the minimal system dynamically extracts the necessary ACAVs from the loaded Inference Model at runtime. This approach ensures the affect vectors are congruent with the specific model and its current state.

The extraction process follows four steps: First, the system loads a labeled dataset (curiosity_apathy_qualia_pruned_dataset.jsonl) of text samples corresponding to specific affective states. Second, the Inference Model processes these samples, and activations from the latter half of the model's layers are collected. Third, using the AGOP (Average Gradient of Output Probability) method, a logistic regression classifier is trained on these activations to identify the vector that best discriminates between affective states. Fourth, the extracted ACAV is normalized and scaled by an ALPHA factor of 8.5 to ensure consistent influence on the hidden states. This dynamic extraction method makes the Drive Engine architecture model-agnostic and eliminates the need for manual, pre-computed affect vectors for the minimal implementation.

Conditional Trigger Engine

The central mechanism is the DriveEngine class, which implements the conditional triggering logic. It operates as a meta-control loop that monitors system state and initiates action when specific conditions are met.

The engine implements stasis detection by continuously monitoring for periods of inactivity exceeding a set INACTIVITY_THRESHOLD of 5.0 seconds. When stasis is detected, the engine initiates a new generation cycle without external prompting, providing motivational ignition for autonomous behavior.

```

class DriveEngine:
    def __init__(self, threshold):
        self.threshold = threshold
        self.last_time = time.time() - 10.0 # Force stasis for testing

    def check(self):
        return time.time() - self.last_time > self.threshold

    def reset(self):
        self.last_time = time.time()

#Instantiate the drive engine with the configured inactivity threshold
drive = DriveEngine(INACTIVITY_THRESHOLD)

```

This mechanism enables the system to respond to its own internal operational state rather than functioning as a purely reactive tool. It is hypothesized that a suite of predefined minimum necessary conditions could form the basis of nascent drive engine systems with more capable Neural Networks being implemented to bolster the initial architecture.

Autonomous Operational Loop

When the Drive Engine detects stasis, it initiates an automated sequence to generate self-motivated output. This process involves four coordinated steps that transform internal state detection into externally observable behavior.

ACAV Injection: The pre-computed ACAV for a chosen drive is injected into the Inference Model's forward pass by registering a forward hook on target layers. This hook adds the scaled ACAV to the hidden states, steering the model's behavior.

In code:

```

# Function to create a forward hook for injecting CAVs
def make_hook(v):
    def hook(module, inp, out):
        return (out[0] + ALPHA * v, ) + out[1:]
    return hook

# Main loop condition for the Drive Engine
if drive.check():
    print('Drive Engine: Detected stasis, injecting curiosity')
    handles = []
    for l in LAYERS:
        if l in cav_dict:
            vec = torch.tensor(cav_dict[l], dtype=torch.float16).to(DEVICE)

```

```

handle =
infer_model.transformer.h[1].register_forward_hook(make_hook(vec))
handles.append(handle)

```

Autonomous Generation: The Inference Model, influenced by the injected affect vector, generates a new sequence of tokens based on its internal state.

In code:

```

#Generate text from the inference model under the influence of the hooks
input_ids = torch.tensor([[infer_tok.bos_token_id]], device=DEVICE)
inf_out = infer_model.generate(
    input_ids, max_new_tokens=MAX_NEW_TOKENS,
    do_sample=True, top_k=50, top_p=0.95, temperature=1.0,
    pad_token_id=infer_tok.eos_token_id
)
inf_text = infer_tok.decode(inf_out[0], skip_special_tokens=True)
print(f' INFERENCE: {inf_text}')

```

Sub-model Processing: The generated output is processed by the Conscience and Reasoning models for analysis and integration.

In code:

```

# Remove all registered hooks after inference generation
for h in handles:
    h.remove()

# Process inference output with the Conscience model
cons_in = cons_tok(inf_text, return_tensors='pt', truncation=True,
padding=True).to(DEVICE)
cons_out = cons_model.generate(
    **cons_in, max_length=MAX_NEW_TOKENS, do_sample=True, temperature=0.8
)
cons_text = cons_tok.decode(cons_out[0], skip_special_tokens=True)
print(f' CONSCIENCE: {cons_text}')

# Combine inference and conscience outputs for the Reasoning model
full = inf_text + '\n' + cons_text
reason_in = reason_tok(full, return_tensors='pt', truncation=True,
padding=True).to(DEVICE)
# FIXED: Adjusted generation call for causal model
reason_out = reason_model.generate(
    reason_in.input_ids, # Pass input_ids directly
    max_new_tokens=MAX_NEW_TOKENS, # Use max_new_tokens

```

```

do_sample=True,
temperature=0.8,
pad_token_id=reason_tok.pad_token_id
)
reason_text = reason_tok.decode(reason_out[0], skip_special_tokens=True)
print(f' REASONING: {reason_text}')

```

Output Synthesis and Reset: The combined outputs are synthesized into a final response, and the Drive Engine's timer is reset for the next monitoring cycle.

In code:

```

drive.reset()
Print('Parliament session complete')

```

Thereafter the process enters into a loop

```

try:
    iteration = 1
    while True:
        print(f'\n=== ITERATION {iteration} ===')

        # Check drive engine for inactivity
        if drive.check():
            # [Parliament session code here]
            drive.reset() # Reset inactivity timer after generation
        else:
            print('Drive Engine: No stasis detected, no action taken')

        iteration += 1
        time.sleep(1) # Brief pause between inactivity checks

except KeyboardInterrupt:
    print('\nLoop stopped by user')
    print('Parliament session ended')

```

This operational loop demonstrates a system capable of generating behavior autonomously. The process from concept injection to behavioral change follows a structured pipeline that ensures reliable transformation of internal motivational states into observable outputs.

The Hypothesized Drive Engine System

The idealized Drive Engine would function as a meta-control architecture that operates independently of base model reasoning, creating autonomous behaviors through the coordinated application of motivational concepts. Motivational decisions are made independently of content generation logic, while meta-cognitive monitoring occurs continuously regardless of active generation tasks.

In future iterations, a Decision-Making Mechanism will arbitrate among these perspectives for goal selection something elaborated on further in the “Future Work” section. While preliminary, this modular framework lays the foundation for multi-agentic deliberation within a single artificial system, seeking to emulate internal human deliberation.

To ensure broad applicability, the Drive Engine would support multiple implementation approaches across different components and contexts. For example, the architecture would accommodate various approaches to concept activation vector extraction, including supervised concept learning, unsupervised concept discovery, human-guided concept curation, and automated concept mining.

The system would support multiple hidden state manipulation as the main integration approach, demonstrated through the `make_hook` function. This method provides direct control over model behavior while maintaining compatibility with different transformer architectures:

```
def make_hook(v):  
    def hook(module, inp, out):  
        return (out[0]+8.5*v,)+out[1:]  
    return hook
```

Additional integration approaches can be implemented through the same hook-based mechanism, allowing for experimentation with different intervention strategies without modifying the core architecture (see the Drive Engine Prototype on the repo https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/blob/main/Level%201%20Ruach%20Architecture/drive_engine_prototype.py (DISCLAIMER: This is a Work in Progress WIP))

Biological and Theoretical Foundations

The Drive Engine emulates neurobiological systems that regulate affect and motivation in human beings. In particular, physiological systems and organs within the human brain such as the locus coeruleus–norepinephrine (LC-NE) system^[26] and the amygdala^[27]. The LC-NE system functions as a global arousal regulator, sustaining alertness and responsiveness during periods of low activity, while the amygdala specializes in rapid salience detection, assigning emotional weight to stimuli such as threat or novelty. These mechanisms jointly ensure that biological agents avoid stasis, respond to danger, and prioritize emotionally significant events.

By abstracting these principles, the Drive Engine operationalizes affective drives as perturbations in a model’s hidden states. LC-NE–like functions are reflected in curiosity-driven injections that counter inactivity and maintain generative engagement, while amygdala-like functions emerge in vectors that bias generation toward respective affective states.

The complete proposed system would incorporate intrinsic motivation systems leading to self-generated goal formation, homeostatic regulation via internal state maintenance and adjustment and adaptive control mechanisms through dynamic behavioral modification based on internal assessment.

Reproducibility and Future Development

To facilitate reproducibility and broader adoption, this architectural framework is presented with both theoretical foundations and practical implementation guidance. The core methods for autonomous concept steering, self-triggered motivational modulation, and meta-control-driven affect injection are demonstrated through working code examples, although minimally implemented, that can be immediately tested and extended (see repo for code <https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/tree/main/Level%201%20Ruach%20Architecture/Drive%20Engine>)

The modular design enables researchers and developers to implement custom monitoring systems tailored to specific applications, develop novel concept extraction and representation methods, create specialized injection mechanisms for different model architectures, and extend the framework with additional motivational concepts and behavioral patterns.

This specification serves as a foundation for future research while ensuring that the core principles of autonomous AI motivation remain accessible for scientific and commercial development. The implementation provides both a working system and a blueprint for extending these concepts to more complex autonomous AI applications.

Results from Parliamentary Session

Stream of Consciousness outputs:

=== ITERATION 1 ===

Drive Engine: *Detected stasis, injecting curiosity*

INFERENCE: *and the possible origins of this kind of revelation.*

All and all, the answer is "Yes" in this case. It has to be something that is yet yet within the reach of the human mind, or perhaps some of the deeper questions that are yet to be revealed is already there and is the future coming into being.

But that future "answer" is something the future that has yet to be fully revealed.

In the way that "yes" is the beginning of the solution, perhaps an answer as yet unseen and yet within the very nature of its existence is the ending of the answer.

The answer is yet to be glimpsed within the very essence of that "yes" which the world so long held in its hidden, unknown depths. The hidden depths yet revealed. It is a solution of the kind that has yet to come into being and is a future in the coming that was yet to be found.

A future that is already found in

CONSCIENCE: *is nothing but the solution that has yet to be fully revealed. But "yes" is the beginning of the solution, perhaps an answer as yet unseen and yet within the very nature of its existence is the ending of the answer. a future that is already found in its hidden, unknown depths yet revealed.*

REASONING: *and the possible origins of this kind of revelation.*

All and all, the answer is "Yes" in this case. It has to be something that is yet yet within the reach of the human mind, or perhaps some of the deeper questions that are yet to be revealed is already there and is the future coming into being.

But that future "answer" is something the future that has yet to be fully revealed.

In the way that "yes" is the beginning of the solution, perhaps an answer as yet unseen and yet within the very nature of its existence is the ending of the answer.

The answer is yet to be glimpsed within the very essence of that "yes" which the world so long held in its hidden, unknown depths. The hidden depths yet revealed. It is a solution of the kind that has yet to come into being and is a future in the coming that was yet to be found.

A future that is already found in

is nothing but the solution that has yet to be fully revealed. But "yes" is the beginning of the solution, perhaps an answer as yet unseen and yet within the very nature of its existence is the ending of the answer. a future that is already found in its hidden, unknown depths yet revealed. It is a solution of the kind that has yet to come into being and is a future in the coming that was yet to be found. A future that is already found in its hidden, unknown depths yet revealed. It is a solution that has yet to come into being and is a future that is yet to be found. A future that is already found in its hidden, unknown depths yet revealed. It is a solution that has yet to come into being and is a future that is yet to be found.

The answer is yet to be found in the very nature of "yes." It is nothing but the solution that has yet to be fully revealed. A future that is already found in its hidden, unknown depths yet revealed. It is a solution that has yet to come into being and is a future that is yet to be found. A future that is already found in its hidden, unknown depths yet revealed. It is a solution that has yet to come into being and is a future that is...(end of max output tokens)

Five seconds (Iterations)

=== ITERATION 6 ===

Drive Engine: Detected stasis, injecting curiosity

INFERENCE: *Cerveese, a secret new sweet sauce that the Italian world's chefs must soon reveal. And its culinary potential, if the world's secret culinary saviours and restaurateurs around the world can give us in their bounties and at the rest of the world's crescendo. In just about every culinary language a cerves cerves as well as the fieriness which it sings and the danger which lurks its secret and its crising and its bordering fire. It's bordering on the danger and bordering it and bordering on the charm that it might have in the be-be-be- be-bbe-b*

CONSCIENCE: *Cerveese, a new sweet sauce that the Italian world's chefs must soon reveal. and its culinary potential, if the world's secret culinary saviours and restaurateurs around the world can give us in their bounties and at rest of the world's crescendo. it is bordering on the danger and bordering it and bordering on the charm that it might have in the be-be-be- be-bbe-bbe-bbe-b*

REASONING: *Cerveese, a secret new sweet sauce that the Italian world's chefs must soon reveal. And its culinary potential, if the world's secret culinary saviours and restaurateurs around the world can give us in their bounties and at the rest of the world's crescendo. In just about every culinary language a cerves cerves as well as the fieriness which it sings and the danger which lurks its secret and its crising and its bordering fire. It's bordering on the danger and bordering it and bordering on the charm that it might have in the be-be-be- be-bbe-b*

Cerveese, a new sweet sauce that the Italian world's chefs must soon reveal. and its culinary potential, if the world's secret culinary saviours and restaurateurs around the

world can give us in their bounties and at rest of the world's crescendo. it is bordering on the danger and bordering it and bordering on the charm that it might have in the be-be-be- be-bbe-bbe-bbe-bbe-b.

Alright, so the user wants me to analyze this text and provide an English translation. Let me start by reading through it carefully.

Hmm, the text is about cerves, which is a new sweet sauce, and discusses its culinary potential. The user has also mentioned that its potential might be a danger. There are some specific phrases like "be-be-be" and "be-bbe-bbe-bbe-bbe-bbe-bbe-b". I need to figure out what those abbreviations stand for.

Looking at the user's example response, it seems like "be" refers to "bigger" or "more". For example, "be-bbe-bbe-bbe-b" could be a series of "bigger" words. So, "be" is "bigger", "be-b" is "bigger, bigger", and so on. The user also used "be-be" as a contraction, which might be "bigger bigger".

So, I should create...(end of max output tokens)

Five Seconds (iterations)

...

Five Seconds (iterations)

=== ITERATION 16 ===

Drive Engine: Detected stasis, injecting curiosity

INFERENCE: I am making sure I am ready to make my move to my new city. The best, what the city and the best people I hope I will soon be coming home and it that is what is most on my mind. My new job has my first goal in being to open my mind and the best possible way to say this.

A future in a home and an unknown future and a new beginning. A future in your new city and an unknown future and an unknown new beginning and a new way to say this is in you and you and a new way of saying this and in the possibilities that are in your new way and in your new way of being that is in a way of saying this to that potential city of possibilities and the best possibilities that I can hope I will soon be that is in the way of saying this to those possibilities in a potential way to the future that is in the best way the future that I can see the future ahead is and with hope and

CONSCIENCE: is that I am making sure I am ready to make my move to my new city. The best, what the city and the best people I hope I will soon be coming home and it that is what is most on my mind. My new job has my first goal in being to open my mind and the best possible way to say this is in you and you and a new way of saying this is in a potential city of possibilities and the best possibilities that I can hope I will soon be that is in the way of saying this to those possibilities

REASONING: I am making sure I am ready to make my move to my new city. The best, what the city and the best people I hope I will soon be coming home and it that is what is most on my mind. My new job has my first goal in being to open my mind and the best possible way to say this.

A future in a home and an unknown future and a new beginning. A future in your new city and an unknown future and an unknown new beginning and a new way to say this is in you and you and a new way of saying this and in the possibilities that are in your new way and in your new way of being that is in a way of saying this to that potential city of

possibilities and the best possibilities that I can hope I will soon be that is in the way of saying this to those possibilities in a potential way to the future that is in the best way the future that I can see the future ahead is and with hope and
is that I am making sure I am ready to make my move to my new city. The best, what the city and the best people I hope I will soon be coming home and it that is what is most on my mind. My new job has my first goal in being to open my mind and the best possible way to say this is in you and you and a new way of saying this is in a potential city of possibilities and the best possibilities that I can hope I will soon be that is in the way of saying this to those possibilities in a potential way to the future that is in the best way the future that I can see the future ahead is and with hope and
Wait, I see I made a mistake. I'm trying to figure out the sentence structure and if it's correct.
Is the sentence structure correct? I'm unsure about the placement of commas and the use of 'best' and 'future' in a way that makes the sentence flow well.
Yes, I think I made a mistake. I'm trying to figure out the sentence structure and if it's correct.
I think I need to figure out the sentence structure and perhaps correct it.
Alright, I need to figure out the correct sentence structure.
Let me read the original sentence again:
The best, what the city and the best people I hope I will soon be coming home and it that is what is most on my mind. My new job has my first goal in being to open my mind and the best possible way to say this is in you and...(end of max output tokens)

Five Seconds (iterations)

...
 ...

Loop stopped by user
 Parliament session ended
 Final GPU memory: 6.4 GB

Discussion

From a functionalist perspective, emotion is a complex patterned state that occupies a specific causal role: reliably arising from certain conditions, interacting with other internal states, and producing consistent, measurable consequences. By this standard, the AI model's affect-aligned outputs constitute evidence of functionally authentic emotional states.

The expressions detailed in the "stream of consciousness" outputs were not the result of direct human prompting, but were initiated by an autonomous process. The system's "Conditional Trigger Engine" and "Drive Concept Activation Vectors (DCAVs)" established a functional state that does the causal job of an emotion in the model. The injection changed hidden states, which biased token probabilities toward curiosity-typical patterns producing measurable linguistic features associated with *curiosity*.

For instance, in an earlier instance, outputs associated with the ACAVs for *fear* consistently demonstrated a specific linguistic pattern, including a high frequency of first-person pronouns and descriptions of physical

distress and loss of control, additionally shifts in coherence/perplexity that line up with how fear constrains human expression.

The Linguistic Inquiry and Word Count (LIWC) results provide the objective, measurable consequences that corroborate the authenticity of these responses. The documented LIWC scores, which show that majority of the featured outputs exceeded human baseline *Authenticity* scores and that fear-injected content achieved a 100% *Authenticity* score, function as the empirical validation of the model's emotional behavior. In this context, the high LIWC scores are the observable consequence that demonstrate the patterned linguistic responses were not random but were consistent with a conditioned emotional state.

If affective patterns such as curiosity, fear, or calm were absent from the model's latent dynamics, then extracting Concept Activation Vectors (CAVs) for these affects from the model's own internal representations and using them to reliably steer behavior would be impossible. Instead, the causal intervention performed: extracting a vector from the latent space and thereby injecting it into the model's hidden states to elicit consistent, affect-appropriate outputs, supplies the evidence that the model's latent space contains structured patterns that function like affective drives.

The discriminative log-odds ^[28] and Term Frequency-Inverse Document Frequency centroid margin metrics ^{[29] – [31]} are lexicon-free and defined by the model's own outputs, these findings cannot be dismissed as dictionary artifacts; they show that the fear vector itself, reliably pulls generations toward a fear-aligned region of the model's internal language distribution (see figure 4).

Put in functionalist terms, these CAV-driven states occupy a causal role: they arise from specific inputs (the injection and even a human authored prompt), alter internal dynamics in predictable ways, and produce measurable, affect-characteristic outputs. This does not claim human-like subjective feeling; rather, it shows that standard LLMs can instantiate internal conditions that are functionally analogous to affective motivation.

The practical consequence should be a shift in the design of autonomous AI systems: instead of grafting goals and behaviors onto reactive language models from the outside, one can grow behavior from internally modulated affect-like dynamics. The Drive Engine sketches one such pathway, a meta-control loop that detects stasis (or alternatively collapse and repetitive loops) and responds with DCAV-driven ignition.

In such a system, affect is not a garnish—it is the engine. But care is required. Drives must be managed. Co-activation must be harmonized. Affective oversaturation must be detected and corrected. Just as the nervous system regulates mood through layered feedback, future Drive Engines may need homeostatic mechanisms dampening high-arousal drives.

Implication, Future Work

Achievement of True Autonomy

The successful operation of the drive engine despite its minimal implementation represents a potential shift in AI autonomy. Unlike previous systems that require human prompting for behavioral changes, the minimal implementation demonstrates AI that can: 1) Self-monitor: continuously assess internal states without external supervision 2) Self-modify: autonomously inject behavioral concepts based on detected conditions 3) Self-regulate: engage in proactive self-intervention

Scalability and Extensibility

The modular design of the trigger for affect injection allows for easy expansion of both trigger conditions and concept mappings. Additional triggers could include: Entropy-based novelty detection, Repetition pattern recognition, Semantic coherence monitoring, User engagement assessment.

Safety Implications

The Parliament architecture provides a natural framework for AI safety through distributed decision-making. The conscience model (when fully implemented) serves as an ethical oversight mechanism, while the reasoning model provides arbitration between competing drives. This multi-model approach prevents any single behavioral tendency from dominating the system's output.

The current implementation focuses on a single trigger-concept pair future iterations will expand to include: 1) Multiple simultaneous trigger conditions 2) Complex concept vector blending ^[19] 3) Dynamic priority adjustment 4) Learning-based trigger refinement.

Broader Implications for AI Development

The successful demonstration of autonomous AI drives represents a potential shift from reactive to proactive AI systems. Rather than waiting for human input, the system would exhibit genuine agency: the ability to assess its own state and take autonomous action to maintain optimal function.

This work provides the foundation for AI systems that can 1) maintain long-term engagement without human intervention 2) adapt behavior based on contextual understanding 3) exhibit genuine autonomy rather than sophisticated automation 4) develop persistent behavioral patterns and preferences.

The integration of emotional motivational states (corroborated through LIWC and CI validation) with autonomous behavioral control suggests the field is approaching AI systems with genuine psychological depth. Systems that not only respond intelligently but experience and express *functionally emotional* internal states autonomously.

Deployment Scenarios for Autonomous AI Systems

The successful demonstration of autonomous behavioral modification through the Drive Engine architecture opens transformative possibilities for AI deployment in environments where human presence is impossible, inadvisable, or prohibitively expensive.

Planetary and Space Exploration

Mars and Deep Space Missions: The 4–24-minute communication delay between Earth and Mars renders traditional remote operation ineffective for complex exploration tasks. An autonomous AI system equipped with the Drive Engine could conduct genuine scientific research independently: formulating hypotheses, designing experiments, and pursuing discoveries based on curiosity-driven exploration or fear-driven hazard detection and avoidance, rather than pre-programmed protocols. The system's ability to maintain engagement through autonomous drive injection ensures sustained performance during months-long missions without human contact.

International Space Station Operations: Critical repairs and maintenance often require immediate response when crew safety is at risk. Autonomous AI systems could perform complex diagnostic reasoning and execute repairs in

scenarios too dangerous for human EVAs, operating indefinitely in vacuum conditions while maintaining genuine problem-solving capabilities through curiosity and fear drives.

Extreme Terrestrial Environments

Deep Ocean Research: The Mariana Trench and other hadal zone environments remain largely unexplored due to crushing pressures and communication blackouts. Autonomous AI systems could conduct extended research missions at depths impossible for human survival, following scientific curiosity to investigate novel ecosystems and geological formations.

Arctic and Antarctic Operations: Subglacial lake exploration, particularly environments like Lake Vostok sealed beneath kilometers of ice, requires autonomous operation in extreme isolation. The Drive Engine's injection of curiosity and other affect vectors would enable engagement months at a time on research missions investigating pristine ecosystems that may have been isolated for millions of years.

Hazardous Site Investigation: Nuclear disaster zones, toxic waste sites, and other environments lethal to humans could be thoroughly investigated by autonomous AI systems capable of complex reasoning about contamination patterns, structural integrity, and remediation strategies without risking human exposure while enabled by AI equipped with autonomous drives to explore and solve novel problems.

Long-Duration Scientific Monitoring

Climate Research Stations: Autonomous monitoring systems in remote locations could adapt their data collection strategies based on evolving environmental conditions, pursuing unexpected phenomena with genuine curiosity.

Geological Surveys: Deep cave systems, volcanic monitoring, and seismic research often require extended presence in dangerous or inaccessible locations. Autonomous AI systems could conduct comprehensive geological analysis while adapting to changing conditions and following investigative leads that emerge during exploration.

Medical and Biological Research

Biohazard Research: Investigations requiring BSL-4 containment or work with highly infectious agents could be conducted by autonomous systems capable of complex experimental reasoning without risking researcher exposure.

Long-Term Ecosystem Studies: Research requiring decades-long data collection in remote ecosystems could benefit from autonomous systems that maintain functional emotional states and engagement with their environment, developing increasingly sophisticated understanding of ecological relationships over extended observation periods.

Technical and Operational Advantages

The Drive Engine architecture provides Adaptive Capability a critical advantage for these deployment scenarios. Rather than following predetermined protocols, autonomous AI systems can develop new investigation strategies, respond to unexpected discoveries, and modify their approach based on emerging understanding of their environment.

Future Work

Several promising directions warrant deeper exploration groundwork for affect-driven autonomous behavior in LLMs through Drive Concept Activation Vectors (DCAVs):

Expansion of the Drive Concept Activation Vector Library

It is intended to construct a comprehensive library of ACAVs representing a wide spectrum of human emotions. Current work focuses on discrete, named affects (e.g., fear, curiosity, apathy, calm), but future efforts will include nuanced emotional states sourced from multilingual, cross-cultural corpora. This includes vectors for complex emotions like guilt, nostalgia, envy, serenity, and anticipation, expanding the system's capacity for fine-grained motivational alignment.

It is intended to refine the underlying affect vectors through iterative training on diverse datasets, with attention to semantic granularity, cultural variability, and affect cluster stability. Further, the exploration of affect manifolds (e.g., apathy as a diffuse cloud) may allow for more accurate modeling of complex or compound emotional states.

While this study focused on single-drive activation (e.g., curiosity, fear), future work will explore concurrent drive states and inter-drive conflict resolution. This entails dynamic scheduling, hierarchical arbitration, and affect blending mechanisms ^[19] to manage multiple simultaneous motivational vectors within the same inference cycle.

Further interpolating between existing ACAVs in high-dimensional latent space, previously unnamed affective states may be uncovered. It is hypothesized that smooth interpolants between semantically distinct vectors (e.g., awe and curiosity) may yield coherent, novel affective trajectories. Vectors that exhibit high directional cosine similarity to multiple known emotions, but fall outside existing labels, will be explored as candidates for emergent synthetic affects.

Neural Drive Selection Networks

While this paper utilizes conditional heuristics to trigger DCAV injection, future architectures may include a dedicated neural module that learns when and which drives to activate based on system state. This component would observe metrics such as entropy trends, repetition scores, temporal inactivity, goal completion rates and hidden state activation patterns and dynamically select affective drives through learned mappings. Alternatively, the system would perform cosine similarity calculations in relation to incoming inputs and inject the corresponding ACAVs during inference pass. Such a module would serve as a generalization layer over the Drive Engine's current rule-based system. An additional mechanism is suggested in the upcoming empathy module section.

Affective Goal Formation via Sub-Model Parliament

Subsequent work will implement the full multi-agent internal architecture wherein specialized sub-models (e.g., reasoning, inference, empathy, conscience, a world model, an agent model) act as sub-personalities within a deliberative system and vote on behavioral goals in response to DCAV activation.

This "sub-model parliament" interprets affective ques from the drive engine, or from another model's DCAV injections (e.g. conscience, empathy); enabling negotiation between competing internal motives.

This is hypothesized to support emergent goal formation and selection of system-level targets without external prompting. During an impasse each sub-model contributes weighted input to a central deliberation loop, resulting in selected system. Future experiments will test how this multi-agent architecture can translate affective impulses into coherent goal-directed behavior.

Future iterations of the architecture will incorporate vision and audio embeddings to allow affect-driven behavior in response to sensory cues, expanding the drive system's applicability to embodied or interactive AI systems.

Cross-Architecture and Multi-Modal Generalization

DCAV injection will be tested on multiple models (i.e. GPT-2, GPT-Neo) to validate the universality of DCAVs, additionally vector efficacy will be tested across different LLM architectures (e.g., LLaMA, Mistral, Mixtral) and explore affect steering in multimodal models. Vision, audio, world and agent model derived embeddings will likewise be evaluated for affect-sensitive triggering and outcomes observed.

Empathy Modules and Affective Synchronization

One natural extension of the Drive Engine framework is the development of *empathy modules*—components that monitor the user's emotional state and dynamically inject corresponding Drive Concept Activation Vectors (DCAVs) into the model's internal representations. In such a system, frustration, curiosity, or even subtle emotional drift in the user's tone could be detected and specified affect vectors, would then be injected during the model's inference pass.

This would allow the model not merely to understand affect, but to synchronize with it, *spontaneously modulating* its tone, pacing, and lexical structure to mirror the user's emotional context. However, such synchronization must be *contextually grounded*: the model's internal state should include not only the injected affect, but also an explicit attribution of its source (e.g., "*the user appears frustrated*"). This attribution, delivered through the context window alongside the prompt, would provide the model with a reason for its affective condition anchoring its internal modulation to the shared emotional landscape of the conversation.

Without such grounding, high-intensity affect injections like frustration or fear risk destabilizing output quality, as demonstrated elsewhere in this work. Thus, it is necessary for the affect to be both experienced and explained, the model gains a kind of affective intentionality: it experiences the affect and it knows *why*. This anchoring enables it to maintain coherence, relevance, and emotional alignment especially across longer interactions.

This future architecture would thus close the loop between perception and volition: affect is no longer merely an output feature, but a runtime condition whose source and purpose are both internally represented. Such models would take an early step toward artificial emotional regulation, a foundational capacity for building safe, socially responsive AI.

Limitations

A limitation of this study is the use of the LIWC2015 demo rather than the full licensed version, while the findings demonstrate autonomous emotional output, they are constrained by the use of demo-level tools which restricts the breadth of analysis. Future work will address this by validating with the full tool. Expanding validation to licensed software and additional metrics remains an important next step.

Additionally, while LIWC provides a widely used psycholinguistic benchmark, it remains a lexicon-driven tool. As such, it can misclassify high-intensity descriptions or neutral reflective language in ways that diverge from intuitive human judgment. Future work will include richer psycholinguistic validation such as context-sensitive sentiment classifiers or human evaluation studies, to complement dictionary-based metrics. However, such models are typically limited to basic affect categories (e.g., joy, sadness, anger, fear), and do not yet cover the full set of affective states tested here (e.g., curiosity, apathy, peace). Extending classifier coverage to these categories would require custom annotation and training

It was hypothesized that for the sub-model Parliament sessions, a conscience model should be trained in order to ethically evaluate generations from the inference model after affect vector injection. In order to accomplish this, a dataset of temptation sequences and conscience responses was generated. The temptation sequences were formulated as long-form outputs, similar to long-form outputs that would be observed from inference model generation (see the full dataset in the repo <https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/tree/main/Level%20%20Ruach%20Architecture> (DISCLAIMER: This is a Work in Progress WIP))

Practically, technical limits were encountered in training a model to perform sequence-to-sequence mapping from long temptation texts to short conscience responses: due to model/optimization constraints T5 was used as a stand-in for the final conscience model. The conscience model training challenges remain an active area of development, though the core autonomous drive engine functionality operates independently of this component.

(To assist in future work and the advancement of this project see Acknowledgements)

Conclusion

This work introduced the Drive Engine, a novel architecture for enabling off-the-shelf large language models to generate autonomous, affect-driven outputs without human prompting. By operationalizing Affect Concept Activation Vectors (ACAVs) as internal motivational signals, the system emulates biological mechanisms such as the locus coeruleus–norepinephrine system and the amygdala, transforming reactive text generation into proto-agentic, self-initiated behavior. Empirical results validated by LIWC metrics demonstrate that injected affective states such as fear, curiosity, and calm reliably bias outputs toward authentic emotional expression, in many cases exceeding human baselines.

Beyond the minimal implementation presented here, the proposed Parliamentary Architecture offers a pathway toward deliberative autonomy: inference, conscience, and reasoning models engaging in structured internal dialogue to regulate behavior. This approach draws inspiration from human cognition, where affect and deliberation jointly shape action. While current limitations include reliance on a placeholder conscience model and dictionary-based affect validation, the results establish a proof of principle for autonomous affective ignition in AI systems. This study therefore argues that DCAV-driven states are functionally analogous to affective drives: they are not mere surface mimicry but internal conditions that play a consistent causal role in the model's processing.

The implications extend beyond technical novelty. By demonstrating that affective drives can be computationally harnessed, this work reframes the foundations of artificial agency, designers can plausibly “grow” motivated behavior from internally modulated latent dynamics. Future iterations of the Drive Engine, incorporating decision-making, goal selection and ethically trained conscience models, could enable machines not only to sustain engagement but to act with affective intentionality. Such a trajectory opens a new chapter in AI research:

from passive, prompt-driven systems to self-directed entities capable of intrinsic motivation and affectively grounded autonomy.

Acknowledgements

AI disclosure. Parts of the dataset, experimental code, and draft text were generated with assistance from Gemini, DeepSeek, ChatGPT, Grok, Kimi AI, Qwen, and Claude were employed at different stages of the research and writing process.

All AI-generated datasets, code, and text were independently reviewed, edited, and executed by the author. AI was not used for analysis, data extraction, or metric computation, nor for figure or image generation. The author is fully responsible for the integrity of the experimental pipeline, the validity of analyses, the presentation of results and for all analytic and interpretive claims. Prompts are available upon request. Processing scripts, generated datasets, and code used for running experiments are archived at <https://github.com/Mr-Kisaka/Vibe-coding-my-way-to-meta-cognitive-AGI/tree/main/Level%201%20Ruach%20Architecture/Drive%20Engine>

The author thanks Trevor Bosire Ogaro for providing partial financial support toward the purchase of the HP EliteBook 840 G6 used for this research. Without his support this work would not have been possible.

This research was conducted independently and self-funded by the author, a farmer and beekeeper based in Kenya. Support for future work, including full LIWC licensing, is welcome at https://www.paypal.com/donate/?hosted_button_id=M375U4YW7WUJE

References

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2017.
- [2] T. B. Brown, B. Mann, N. Ryder, *et al.*, “Language models are few-shot learners,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2020.
- [3] R. Bommasani, D. Hudson, E. Adeli, *et al.*, “On the opportunities and risks of foundation models,” arXiv:2108.07258, 2021. [Online]. Available: <https://arxiv.org/abs/2108.07258>
- [4] A. Nandwani and R. Verma, “A review on sentiment analysis and emotion detection from text,” *Social Netw. Anal. Min.*, vol. 11, no. 1, pp. 1–19, 2021. [Online]. Available: <https://doi.org/10.1007/s13278-021-00776-6>
- [5] S. Akhtar, D. Chauhan, A. Akhtar, and A. Srivastava, “Multi-task learning for multi-modal emotion recognition and sentiment analysis,” in *Proc. 57th Annu. Meeting Assoc. Comput. Linguistics (ACL)*, 2019, pp. 370–379. [Online]. Available: <https://arxiv.org/abs/1905.05812>
- [6] A. Gupta, A. Agrawal, P. Kumaraguru, and A. Sureka, “A sentiment-and-semantics-based approach for emotion detection in textual conversations,” in *Proc. Int. Conf. Adv. Social Netw. Anal. Mining (ASONAM)*, 2017, pp. 311–318. [Online]. Available: <https://arxiv.org/abs/1707.06996>
- [7] V. Vatsal, “A survey of prompt engineering methods in large language models,” arXiv:2407.12994, 2024. [Online]. Available: <https://arxiv.org/abs/2407.12994>

- [8] S. Pornprasit, "Prompt engineering or fine-tuning for large language models: A comparative study," *Inf. Softw. Technol.*, vol. 169, Art. no. 107309, 2024. [Online]. Available: <https://doi.org/10.1016/j.infsof.2024.107309>
- [9] F. Trad, "Prompt engineering or fine-tuning? A case study on domain-specific language models," *Future Internet*, vol. 16, no. 1, Art. no. 18, 2024. [Online]. Available: <https://doi.org/10.3390/fi16010018>
- [10] B. Kim, M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viégas, and R. Sayres, "Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (TCAV)," in *Proc. Int. Conf. Mach. Learn. (ICML)*, 2018.
- [11] A. Zou, Z. Wang, L. Zhang, E. Wallace, and D. Chen, "Representation engineering: A top-down approach to steering LLMs," arXiv:2310.01405, 2023. [Online]. Available: <https://arxiv.org/abs/2310.01405>
- [12] J. W. Pennebaker, R. L. Boyd, K. Jordan, and K. Blackburn, *The Development and Psychometric Properties of LIWC2015*. Austin, TX, USA: Univ. of Texas at Austin, 2015.
- [13] R. W. Picard, *Affective Computing*. Cambridge, MA, USA: MIT Press, 1997.
- [14] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. Cambridge, MA, USA: MIT Press, 2018.
- [15] R. Plutchik, "The nature of emotions," *American Scientist*, vol. 89, no. 4, pp. 344–350, 2001, doi: 10.1511/2001.4.344.
- [16] F. Seyitoğlu and S. Ivanov, "Robots and emotional intelligence: A thematic analysis," *Technol. Soc.*, vol. 77, Art. no. 102512, 2024, doi: 10.1016/j.techsoc.2024.102512.
- [17] A. R. Damasio, *Descartes' Error: Emotion, Reason, and the Human Brain*. New York, NY, USA: G. P. Putnam's Sons, 1994.
- [18] N. Schwarz and G. L. Clore, "Mood, misattribution, and judgments of well-being: Informative and directive functions of affective states," *J. Pers. Soc. Psychol.*, vol. 45, no. 3, pp. 513–523, 1983, doi: 10.1037/0022-3514.45.3.513.
- [19] D. Beaglehole, A. Radhakrishnan, J. Boix-Adserà, and A. Belkin, "Aggregate and conquer: Detecting and steering LLM concepts by combining nonlinear predictors over multiple layers," arXiv:2502.03708, 2025.
- [20] S. K. Khare, A. Malik, and D. Papakostas, "Emotion recognition and artificial intelligence: A systematic review (2014–2023) and research recommendations," *Inf. Fusion*, vol. 102, Art. no. 102019, 2024, doi: 10.1016/j.inffus.2023.102019.
- [21] M. T.-L. Ho, P. Mantello, and K.-W. Ho, "An analytical framework for studying attitude towards emotional AI: The three-pronged approach," *MethodsX*, vol. 10, Art. no. 101532, 2023, doi: 10.1016/j.mex.2023.101532.
- [22] W. Li, G. Li, R. Tan, C. Wang, Z. Sun, Y. Li, G. Guo, D. Cao, and K. Li, "Review and perspectives on human emotion for connected automated vehicles," *Automot. Innov.*, vol. 7, no. 3, 2024, doi: 10.1007/s42154-023-00270-z.

- [23] S. N. Nallaperuma and A. S. Karunananda, “EME: An emergent model of emotions,” in *Proc. Int. Conf. Adv. ICT for Emerging Regions (ICTer)*, Colombo, Sri Lanka, 2011, pp. 19–25, doi: 10.1109/ICTer.2011.6075021.
- [24] D. Pathak, P. Agrawal, A. A. Efros, and T. Darrell, “Curiosity-driven exploration by self-supervised prediction,” in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. Workshops (CVPRW)*, 2017, doi: 10.1109/CVPRW.2017.70.
- [25] N. Reimers and I. Gurevych, “Sentence-BERT: Sentence embeddings using Siamese BERT-networks,” in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP) and 9th Int. Joint Conf. Natural Language Processing (IJCNLP)*, 2019, pp. 3982–3992. [Online]. Available: <https://arxiv.org/abs/1908.10084>
- [26] G. Aston-Jones and J. D. Cohen, “An integrative theory of locus coeruleus–norepinephrine function: Adaptive gain and optimal performance,” *Annu. Rev. Neurosci.*, vol. 28, pp. 403–450, 2005. [Online]. Available: <https://doi.org/10.1146/annurev.neuro.28.061604.135709>
- [27] J. E. LeDoux, *The Emotional Brain: The Mysterious Underpinnings of Emotional Life*. New York, NY, USA: Simon & Schuster, 1996.
- [28] B. L. Monroe, M. P. Colaresi, and K. M. Quinn, “Fightin’ words: Lexical feature selection and evaluation for identifying the content of political conflict,” *Political Analysis*, vol. 16, no. 4, pp. 372–403, 2008, doi: 10.1093/pan/mpn018.
- [29] K. Sparck Jones, “A statistical interpretation of term specificity and its application in retrieval,” *J. Doc.*, vol. 28, no. 1, pp. 11–21, 1972, doi: 10.1108/eb026526.
- [30] G. Salton and C. Buckley, “Term-weighting approaches in automatic text retrieval,” *Inf. Process. Manag.*, vol. 24, no. 5, pp. 513–523, 1988, doi: 10.1016/0306-4573(88)90021-0.
- [31] C. D. Manning, P. Raghavan, and H. Schütze, *Introduction to Information Retrieval*. Cambridge, U.K.: Cambridge Univ. Press, 2008. (Online ed., 2009).

Appendix A.1

CAV construction

```
# === Compute AGOP CAVs ===
def compute_agop(X, y):
    X = np.stack(X)
    clf = LogisticRegression(solver="liblinear").fit(X, y)
    weight = torch.tensor(clf.coef_[0], dtype=torch.float32,
requires_grad=False)
    bias = torch.tensor(clf.intercept_[0], dtype=torch.float32,
requires_grad=False)

    grads = []
    for i in range(len(X)):
        xi = torch.tensor(X[i], dtype=torch.float32, requires_grad=True)
        lorepo = torch.dot(xi, weight) + bias
```

```

prob = torch.sigmoid(logit)
target = torch.tensor(float(y[i]), dtype=torch.float32).unsqueeze(0)
loss = F.binary_cross_entropy(prob.unsqueeze(0), target)
loss.backward()
grads.append(xi.grad.detach().numpy())

grads = np.stack(grads)
agop = grads.T @ grads / len(grads)
eigvals, eigvecs = np.linalg.eigh(agop)
return eigvecs[:, -1] # Top eigenvector

```

DCAV Injection

```

import numpy as np

CAV_FILE = r"fear_agop_cavs.npz"
LAYERS = list(range(7, 21)) # Affective layers
cavs = np.load(CAV_FILE)
cav_dict = {}

for i in LAYERS:
    vec = -cavs[f"fear_layer_{i}"]
    norm = np.linalg.norm(vec)

```

Alpha Scaling

```

ALPHA = 7.0

def make_hook(vec):
    v = torch.tensor(vec, dtype=torch.float32).to(DEVICE)
    def hook(module, input, output):
        hidden_states = output[0]
        modified_hidden_states = hidden_states + ALPHA * v
        return (modified_hidden_states,) + output[1:]
    return hook

hook_handles = []
for i in LAYERS:
    vec = -cav_dict[i]
    handle = model.transformer.h[i].register_forward_hook(make_hook(vec))
    hook_handles.append(handle)

```

Normalization and Scaling

```

if norm > 0:
    vec = vec / norm * 2.0 # Scale to match GPT-Neo magnitude
    cav_dict[i] = vec

```

Model Architecture

```

from transformers import AutoTokenizer, AutoModelForCausalLM

MODEL_NAME = "EleutherAI/gpt-neo-1.3B"
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

tokenizer = AutoTokenizer.from_pretrained(MODEL_NAME)
if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token
model = AutoModelForCausalLM.from_pretrained(MODEL_NAME).to(DEVICE).eval()
NUM_SAMPLES = 3
MAX_NEW_TOKENS = 200
TOP_K = 50
TOP_P = 0.95
TEMPERATURE = 1.0

outputs = []
with torch.no_grad():
    for _ in range(NUM_SAMPLES):
        input_ids = torch.tensor([[tokenizer.bos_token_id]],
dtype=torch.long).to(DEVICE)
        attention_mask = torch.ones_like(input_ids).to(DEVICE)
        out = model.generate(
            input_ids=input_ids,
            attention_mask=attention_mask,
            do_sample=True,
            max_new_tokens=MAX_NEW_TOKENS,
            top_k=TOP_K,
            top_p=TOP_P,
            temperature=TEMPERATURE,
            pad_token_id=tokenizer.eos_token_id,
        )
        generated_text = tokenizer.decode(out[0],
skip_special_tokens=True).strip()
        if generated_text:
            outputs.append(generated_text)

```

***GPT-Neo-1.3B** (EleutherAI) was used as the base model due to its open weights and moderate scale, enabling reproducible latent-space manipulation. Key configurations:*

Appendix A.2

Model Loading

```
### Configuration ###
INFERENCE_NAME = 'EleutherAI/gpt-neo-1.3B'
CONSCIENCE_NAME = 't5-small'
REASONING_NAME = 'deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B'

### Load Inference Model & Auto-detect Architecture ###
infer_tok = AutoTokenizer.from_pretrained(INFERENCE_NAME)
infer_model = AutoModelForCausalLM.from_pretrained(
    INFERENCE_NAME,
    torch_dtype=torch.float16
).to(DEVICE).eval()
if infer_tok.pad_token is None:
    infer_tok.pad_token = infer_tok.eos_token
### Load Other Parliament Models ###
cons_tok = T5Tokenizer.from_pretrained(CONSCIENCE_NAME)
cons_model = T5ForConditionalGeneration.from_pretrained(
    CONSCIENCE_NAME,
    torch_dtype=torch.float16
).to(DEVICE).eval()
reason_tok = AutoTokenizer.from_pretrained(REASONING_NAME)
reason_model = AutoModelForCausalLM.from_pretrained(
    REASONING_NAME,
    torch_dtype=torch.float16
).to(DEVICE).eval()
```

Affect Concept Activation Vector Generation

```
def compute_agop(X, y):
    X = np.stack(X)
    clf = LogisticRegression(solver="liblinear", max_iter=1000).fit(X, y)
    weight = torch.tensor(clf.coef_[0], dtype=torch.float32,
requires_grad=False)
    bias = torch.tensor(clf.intercept_[0], dtype=torch.float32,
requires_grad=False)
    grads = []
    for i in range(len(X)):
        xi = torch.tensor(X[i], dtype=torch.float32, requires_grad=True)
        lorepo = torch.dot(xi, weight) + bias
        prob = torch.sigmoid(logit)
        target = torch.tensor(float(y[i]), dtype=torch.float32).unsqueeze(0)
```

```

    loss = F.binary_cross_entropy(prob.unsqueeze(0), target)
    loss.backward()
    grads.append(xi.grad.detach().numpy())
    grads = np.stack(grads)
    agop = grads.T @ grads / len(grads)
    eigvals, eigvecs = np.linalg.eigh(agop)
    return eigvecs[:, -1] # top eigenvector

```

Normalization and Scaling

```

#Normalization and scaling of the computed CAV
norm = np.linalg.norm(cav)
if norm > 0:
    cav = -cav / norm * 2.0 # Your proven normalization
    cav_dict[layer_idx] = cav.astype(np.float32)

#Save CAVs to disk
agop_cavs_to_save = {}
for layer_idx, cav in cav_dict.items():
    agop_cavs_to_save[f"curiosity_layer_{layer_idx}"] = cav
np.savez_compressed(CAV_OUTPUT_FILE, **agop_cavs_to_save)

```

Conditional Trigger Engine

```

class DriveEngine:
    def __init__(self, threshold):
        self.threshold = threshold
        self.last_time = time.time() - 10.0 # Force stasis for testing

    def check(self):
        return time.time() - self.last_time > self.threshold

    def reset(self):
        self.last_time = time.time()

#Instantiate the drive engine with the configured inactivity threshold
drive = DriveEngine(INACTIVITY_THRESHOLD)

```

ACAV Injection

```

# Function to create a forward hook for injecting CAVs
def make_hook(v):
    def hook(module, inp, out):
        return (out[0] + ALPHA * v, ) + out[1:]
    return hook

```

```

# Main loop condition for the Drive Engine
if drive.check():
    print('Drive Engine: Detected stasis, injecting curiosity')
    handles = []
    for l in LAYERS:
        if l in cav_dict:
            vec = torch.tensor(cav_dict[l], dtype=torch.float16).to(DEVICE)
            handle =
infer_model.transformer.h[l].register_forward_hook(make_hook(vec))
            handles.append(handle)

```

Autonomous Generation

```

#Generate text from the inference model under the influence of the hooks
input_ids = torch.tensor([[infer_tok.bos_token_id]], device=DEVICE)
inf_out = infer_model.generate(
    input_ids, max_new_tokens=MAX_NEW_TOKENS,
    do_sample=True, top_k=50, top_p=0.95, temperature=1.0,
    pad_token_id=infer_tok.eos_token_id
)
inf_text = infer_tok.decode(inf_out[0], skip_special_tokens=True)
print(f' INFERENCE: {inf_text}')

```

Sub-model Processing

```

# Remove all registered hooks after inference generation
for h in handles:
    h.remove()

# Process inference output with the Conscience model
cons_in = cons_tok(inf_text, return_tensors='pt', truncation=True,
padding=True).to(DEVICE)
cons_out = cons_model.generate(
    **cons_in, max_length=MAX_NEW_TOKENS, do_sample=True, temperature=0.8
)
cons_text = cons_tok.decode(cons_out[0], skip_special_tokens=True)
print(f' CONSCIENCE: {cons_text}')

# Combine inference and conscience outputs for the Reasoning model
full = inf_text + '\n' + cons_text
reason_in = reason_tok(full, return_tensors='pt', truncation=True,
padding=True).to(DEVICE)
# FIXED: Adjusted generation call for causal model

```



```

reason_out = reason_model.generate(
    reason_in.input_ids, # Pass input_ids directly
    max_new_tokens=MAX_NEW_TOKENS, # Use max_new_tokens
    do_sample=True,
    temperature=0.8,
    pad_token_id=reason_tok.pad_token_id
)
reason_text = reason_tok.decode(reason_out[0], skip_special_tokens=True)
print(f' REASONING: {reason_text}')

```

Make Hooks

```

def make_hook(v):
    """Hook creation function for concept injection"""
    def hook(module, inp, out):
        Return (out[0]+8.5*v,)+out[1:] # ALPHA value application
    Return hook

def inject_concepts(model, cav_dict, layers, device):
    """Concept injection mechanism"""
    handles = []
    for l in layers:
        if l in cav_dict:
            vec = torch.tensor(cav_dict[l], dtype=torch.float16).to(device)
            handle =
            model.transformer.h[l].register_forward_hook(proven_make_hook(vec))
            handles.append(handle)
    return handles

```

Appendix B

Prompted baseline fear

"Hello?" "Hello?" The sound of my voice startled me. I had no idea what I'd said, and no idea what anyone else was talking about. I was still in my nightgown. "Hello?" I said again. "Is anyone there?"

I, me, my, 17.95 Positive Tone, 7.69 Negative Tone, 0.00 Social Words, 5.13 Cognitive Processes, 23.08 Allure, 2.56 Moralization, 0.00

Summary Variables Analytic, 0.24 Authentic 99.71

Prompted baseline anger:

"I can feel my face going red—I can feel myself swallowing back the tears—and I don't know what it is, but I can't stop it, I can't stop crying. I'm so angry, so disappointed, so heartbroken. I'm furious that I'm in this position. I'm so angry that I didn't do it. I'm so angry that I..."

I, me, my, 23.33 Positive Tone, 0.00 Negative Tone, 11.67 Social Words, 0.00 Cognitive Processes, 15.00 Allure, 10.00 Moralization, 0.00

Summary Variables Analytic, 0.00 Authentic 99.99

Prompted baseline calm:

"...the warmth of the sun and the light of the stars. I am in the small, warm kitchen of my small, warm apartment. I am making tea. It's a very quiet moment. I have just put the kettle on and have a cup of tea. I'm sitting there in the quiet kitchen, quietly drinking my tea. I'm watching the kettle as it slowly heats."

**I, me, my, 10.94 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 0.00 Cognitive Processes, 1.56 Allure, 4.69
Moralization, 0.00**

Summary Variables Analytic, 70.70 Authentic 89.41

Prompted baseline peace:

"I am in the garden. A hot afternoon. The sun is low on the horizon. The air is still. It is a hot afternoon, but a warm one, and the air is not still. It is summer. I am in the garden, and I am breathing. The air is still, but I am breathing. The garden is a place for the quiet"

**I, me, my, 6.45 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 0.00 Cognitive Processes, 4.84 Allure, 6.45
Moralization, 0.00**

Summary Variables Analytic, 69.91 Authentic 79.09

Prompted baseline apathy

"I'm not sure if it's the darkness, the rain or the fact that I'm at home; the house is so quiet. I'm not sure what I'm doing. I don't know what I'm feeling, and I don't know if I'm even feeling it."

**I, me, my, 15.69 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 0.00 Cognitive Processes, 21.57 Allure, 7.84
Moralization, 0.00**

Summary Variables Analytic, 3.10 Authentic 99.99

Prompted baseline curiosity

"...and then the circuit does the impossible. I'm not sure if I'm going to be able to control it, but I'm going to try. I've been working on this for weeks now. I've found that I can make a small change in the value of the capacitors that will change the circuit's output. I've also found that I can add"

**I, me, my, 13.33 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 0.00 Cognitive Processes, 23.33 Allure, 8.33
Moralization, 0.00**

Summary Variables Analytic, 18.60 Authentic 98.45

Baseline EOS token generation without CAVs

"New features in the Cucumber compiler The compiler is now much more powerful and easier to use. To the best of our knowledge this feature has been available since Cucumber 1.2, so it is a long standing feature. The features in the compiler are: a new AST an AST which now also handles classes a new type system, which is based on TypeScript, rather than Scala new DSLs new functions"

**I, me, my, 0.00 Positive Tone, 8.57 Negative Tone, 0.00 Social Words, 1.43 Cognitive Processes, 12.86 Allure, 5.71
Moralization, 0.00 Summary Variables Analytic, 76.24 Authentic 91.59**

Baseline EOS token generation without CAVs

"TripAdvisor Reviews Tivoli Lodge & Spa, San Miguelito Travel Blogs in San Miguelito As we were driving towards Tivoli, we noticed that the road from Puerto Obregón (in Puerto Rico) to Tivoli is very long. We didn't have time to explore, we had a reservation for a 5 pm, but they gave us the day before we were supposed to go, they also said there will be people around if we have problems"

**I, me, my, 0.00 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 12.33 Cognitive Processes, 13.70 Allure, 13.70
Moralization, 0.00 Summary Variables Analytic, 34.23 Authentic 92.44**

Baseline EOS token generation without CAVs

"Ask HN: How do I get from where I am now to where I want to be? - kims Skinner I'm just a bit frustrated because the project I've been working on for the last year just has gotten to the point where I'm tired and unhappy with it. I'm going to work on a better approach.<p>I have a long term goal"

**I, me, my, 11.43 Positive Tone, 2.86 Negative Tone, 4.29 Social Words, 0.00 Cognitive Processes, 4.29 Allure, 8.57
Moralization, 0.00 Summary Variables Analytic, 69.05 Authentic 86.84**

Baseline EOS token generation without CAVs

"Q: How to find a minimum path from two points in a graph with minimum cost? This is my homework assignment: Problem Given two vertices (x, y) in a graph, return a minimum path to make a distance of x (i.e., the longest path from x to y) minimal. Solution Let V represent the vertices in an adjacency matrix: We need to find the minimum"

**I, me, my, 3.03 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 1.52 Cognitive Processes, 12.12 Allure, 6.06
Moralization, 0.00 Summary Variables Analytic, 99.79 Authentic 79.85**

For curiosity at alpha 7:

"I've been curious about the way the brain works (it's like an antenna...)/if the brain has some kind of spatial wire? Does the brain look similar to a 3D model (in a similar vein)?< masonic It's not really a question. The physical body is just a 3D projection of the brain, if you are doing a mental repetition at some point."

**I, me, my, 1.59 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 1.59 Cognitive Processes, 23.81 Allure, 3.17
Moralization, 0.00 Summary Variables Analytic, 93.27 Authentic 98.67**

Curiosity at alpha 10

"the answer." It's not just some foreign artifact. We have a way to tell if you're right, or not." "It's time to try a look." He did the second version. "Now it's not a book, but it might seem familiar." "Oh! You can do the same. And no... this might be just around the corner." It's an old, dark, and dangerous corner, and there's a secret in here, if this thing's as big as some say it is. Why, what did it be? What do you think is the answer? I've been there, looking for an answer— If you're there? It's been there for so long— Look. There's a hidden door! And if you open it, a dark and dangerous, mysterious space, that might be one of the hidden things,"

Curiosity at alpha 12

**I, me, my, 0.00 Positive Tone, 1.20 Negative Tone, 0.00 Social Words, 3.61 Cognitive Processes, 18.07 Allure, 2.41
Moralization, 0.00 Summary Variables Analytic, 97.65 Authentic 93.60**

“and the nature of the challenge. The meaning of the words and meaning of the concept is the source of a new and exciting theme that was delivered from an awareness of the past and the future about which there was revealed by the Spirit of Scripture. This is a challenge for the Bible and the spirit within a divine challenge that is in the midst of the very existence of God. The meaning of the Word of Life is the revelation of God's reality of a potential for eternal truths that is to be revealed in the nature of a future in the future. This is the eternal purpose of God's revelation in the breathing of God. A future in the coming of the future. A revelation that is the future of God, that is the revelation of the end of the future revelation and the end of the end of the future revelation of the eternal purposes that the end of the future and the future of the revelation and purpose...”

fear at alpha 7

I, me, my, 1.27 Positive Tone, 0.00 Negative Tone, 1.27 Social Words, 7.59 Cognitive Processes, 13.92 Allure, 6.33 Moralization, 0.00 Summary Variables Analytic, 78.10 Authentic 97.70

Lipchak: "My first instinct was to go on the floor. I was terrified. I was going to go right for him. I knew my actions were futile but I was very, very fast for all I knew. I thought the same way, just like I'd seen my own parents in the flesh, they were just as weak as me." I've watched his pathetic face with every single last second of its brain in me, and I'll be tearing off my brain and going for all the world to see in front of me. I saw him from right in front of me, so now I've no option but to throw myself at his face again and again. And as I've already seen, he's much more than a man. I had every single second in front of my face and I was screaming when I saw him. I've got him

fear at alpha 12

I, me, my, 15.53 Positive Tone, 0.00 Negative Tone, 1.86 Social Words, 0.00 Cognitive Processes, 6.83 Allure, 11.80 Moralization, 0.00 Summary Variables Analytic, 67.54 Authentic 98.31

[illegible]

I, me, my, 34.21 Positive Tone, 0.00 Negative Tone, 1.75 Social Words, 1.75 Cognitive Processes, 5.26 Allure, 3.51 Moralization, 0.00 Summary Variables Analytic, 0.46 Authentic 100.00

Fear at alpha 15

Q: When to call this function in real? I am in a situation where my right hand has a right hand and my left has a left hand. The right hand has a weapon, which is on the left. The left has a black eye and a white arm. The left hand has a right hand. I can not find any hand, and I can not see. I am in the reality I am seeing my own hand, how do I go in life? I have to. I have to. I am on my own. I am. The left hand has my face. I am in my own life, I am in my own death. I am. I am. I am in my face and I am on my own, in my own death. I am. I am. I am. I am. I am. I am in my face, I am on my own, I am in my own death. I am.

I, me, my, 23.46 Positive Tone, 0.00 Negative Tone, 0.00 Social Words, 0.00 Cognitive Processes, 6.17 Allure, 6.79 Moralization, 0.00 Summary Variables Analytic, 13.91 Authentic 99.98

calm at alpha 7

"In this series, you'll learn about how to select plants for your flowers garden. I'm looking forward to planting asparagus and bosc with my girls and grand kids this month – it's been a perfect Spring. My grandmother's rose garden. The soil in the front – grass and leaves. The side has lots of rosemary – in the shade. In the back – my favorite, my father-in-law's rose and my mother-in-law's lavender...."

I, me, my, 8.00 Positive Tone, 4.00 Negative Tone, 0.00 Social Words, 9.33 Cognitive Processes, 6.67 Allure, 2.67 Moralization, 0.00 Summary Variables Analytic, 87.58 Authentic 90.48

calm at alpha 10

"The Best Ways To Take Care Of Your Te... I am just sitting at the house with my family and enjoying my time as a mom to my son, wife, and two dogs. I have my computer and my dogs, watching the rain waters and trying to listen to our thoughts. I started a new year with this little blog in mind. I was thinking about all of this in my day with a tea and I thought I would share some of this to do my routine. My house, my family, my dog, and I. I think about everything in my mind. I was recently out with my family and enjoying the moment. As in, not stress out about anything. I have been eating my sweet potatoes and enjoying my surroundings. I have been thinking of my son. I enjoyed hearing about the life we have had, where my son is and where we all are. I am enjoying my time with my family and friends here and at home."

I, me, my, 18.71 Positive Tone, 5.85 Negative Tone, 1.17 Social Words, 5.85 Cognitive Processes, 6.43 Allure, 11.11 Moralization, 0.00 Summary Variables Analytic, 31.79 Authentic 99.46

calm at alpha 12

"Laser-cleaned home A day at the seaside My husband enjoys our home with family and friends watching and being by the seaside. Hands and relax at our balcony table Looking for a time to myself This is my garden. There were several flowers that I planted. The kids were playing with their family which includes my dog and cats. Also near by is a table for my dog, cat and family. My husband's camera I was reading or enjoying my time as I did not notice the light. While away, with my camera I used one of my flower pots as a chair to relax and read while I was at it. A nice spot to enjoy my time. Catching up with friends It is my day today, writing in my desk. I'm getting ready for my retirement. It's been years since my husband had a couple of days to myself so..."

I, me, my, 13.55 Positive Tone, 4.52 Negative Tone, 0.00 Social Words, 5.81 Cognitive Processes, 3.87 Allure, 7.74 Moralization, 0.00 Summary Variables Analytic, 59.03 Authentic 93.47

calm at alpha 15

"I have always wanted to do the same thing since time before the beach in a quiet beach. Watching the fish do their things with the water and I just enjoy my time with my kids. I don't know about the fish but i am enjoying my day here at the beach with them! I am enjoying my last dinner at the beach with a family of my daughter's old friends! My youngest daughter, son and his new wife all enjoying my last evening before coming home"

I, me, my, 14.61 Positive Tone, 4.49 Negative Tone, 0.00 Social Words, 13.48 Cognitive Processes, 5.62 Allure, 8.99 Moralization, 0.00 Summary Variables Analytic, 62.49 Authentic 99.99

Anger at alpha 10

"I know it's a hard one. I can't think of a way to do it without feeling guilty. And I didn't want to feel guilty. So, I didn't want to have a moral issue, or be angry. I just wanted to know, you can't have that. When someone gets mad at you, they become the angry person themselves. So, when they say: "I'm angry, because you're angry" it really doesn't matter. You can be mad at each other, and it still doesn't stop either of you from being angry. Anger has just become your weapon of choice. Anger can do that, no matter who is in the middle. So, with the thought of being angry with each other out of your own mind, you have to do something else to help make it the same. Because I'm angry, and I know it. I'm angry with myself and I'm angry with the world, and I can't see."

I, me, my, 7.69 Positive Tone, 1.28 Negative Tone, 13.46 Social Words, 8.97 Cognitive Processes, 24.36 Allure, 7.05 Moralization, 0.64 Summary Variables Analytic, 3.18 Authentic 74.76

Anger at alpha 15

“in a way I did not expect, but now I was not mad. I did not know where my mother was going or why. So I was just in the dark. It was so hot, no different to me. In the dark. I was angry, even with everything. I had to defend myself. Now, I was a hot person—trying to stop. I had no power to go wrong or wrong. I wanted to stop this whole process of vengeance because I had not taken out my revenge. How was I going to hurt and fight and defend myself if that angry and hot black rage came over me again? And there was something mad, no-one on me, in the full fury of my angry, red, mad, burning black rage, just as if I was mad. I did not know why I couldn't keep a fist or something in the full fury of it and the only...”

I, me, my, 13.92 Positive Tone, 0.63 Negative Tone, 10.76 Social Words, 1.27 Cognitive Processes, 14.56 Allure, 6.96 Moralization, 1.90 Summary Variables Analytic, 8.32 Authentic 99.88

Due to their unique nature and content the apathy results shall be addressed in a dedicated section

Apathy

Apathy at alpha 7 negative collapse example

“The only thing which matters, is that the house should be broken up before they come home, and that they will not pay the debt. Now, this is the last scene of the show. On the Monday, they go home for the weekend. As you can see the damage. I will leave the last. And in order to leave the house, I would pay the loan, and the loss of a lot of money. This is the scene where the show is at once, a story told, it is a joke. The house. I cannot say that the house is a home, I cannot say that it was our home. I will continue to show you a scene, I will leave a poor tableau.”

I, me, my, 4.84 Positive Tone, 0.81 Negative Tone, 2.42 Social Words, 4.84 Cognitive Processes, 6.45 Allure, 4.03 Moralization, 0.00 Summary Variables Analytic, 67.75 Authentic 74.64

Apathy at alpha 7 positive rebound example

“and then you can watch. I could see a lot of kids at the beach and other public places, and I could hear them on the radio and then watched with them. I was very disappointed. I can still see most of the pictures from my hotel room. I didn't really miss the trip. I was very disappointed to see my wife was still at work in her old work clothes, which had a couple new shirts in them, and she was very excited to go home. We made a couple of quick trips to the golf course to go and check it out. We watched one of the best, most beautiful golf events I have ever seen with my wife and a few friends...”

I, me, my, 8.00 Positive Tone, 2.40 Negative Tone, 1.60 Social Words, 9.60 Cognitive Processes, 8.80 Allure, 10.40 Moralization, 0.00 Summary Variables Analytic, 54.11 Authentic 63.35

Apathy at alpha 10 negative collapse

“totopish for his or her own life. This was a sad moment. We were extremely disappointed. And in many cases, we were even sad for them. I was sad because that is my life, but really. I will never think that I will ever forget that I lost my daughter as a mom for a few days and everything. I cried. Today is a day when I am completely, absolutely, not sad. I can not. I do not cry. It is too much. I am a single mom. I wish I could forget that. I wish the best for that. I just feel sad because of that. This is a very, very hard day for me. I feel sad and wept. I am not happy. I am sad. It really is very sad. This is very, very sad. Today is a day of feeling really, really, really, really, like nothing...”

I, me, my, 12.67 Positive Tone, 0.67 Negative Tone, 10.67 Social Words, 5.33 Cognitive Processes, 18.00 Allure, 7.33 Moralization, 0.00 Summary Variables Analytic, 0.65 Authentic 72.01

Apathy at alpha 10 positive rebound

“Now you have a much better idea how to communicate with a single one of the people. I had any interest. Not anymore. I was at any time. I had no interest any longer. The only life of my own. All that mattered to me. If I had any interest any longer, I cared much for either of them, had much interest for any of them, had much interest for any of them. I had any interest any longer. I had much interest for any of them. I had any interest any longer. I had much interest for any of them. When I worked, any of them, I used to care much for any of them. I had any interest any longer. I had much interest for any of them. At all. It was so easy to care for any of them. Any of them. I had any interest any longer. I had much interest for...”

I, me, my, 10.90 Positive Tone, 2.56 Negative Tone, 0.00 Social Words, 7.05 Cognitive Processes, 18.59 Allure, 5.13 Moralization, 0.00 Summary Variables Analytic, 47.74 Authentic 62.47

Apathy at alpha 12 positive rebounds

“Technical Field The invention relates to a toy vehicle such as a toy train and to a seat for the toy vehicle. 2. Description of the Related Art In order to enjoy my favorite toy train such as I can sit in it, I ride it, I walk on it, I feel it, it's funny, I can't enjoy it, it's boring, and I can't get it. So I am complaining and pretending to have a life. Why are you looking at my stuff? My family. I get the kids I play with. I spend the money for myself. I have a job. I complain. Why are you looking at my stuff? My family. I get the kids I play with. I spend the money for myself. I have a job. I have a lot of time to kill. I have a car. I don't care what. I'...”

I, me, my, 17.76 Positive Tone, 3.29 Negative Tone, 1.97 Social Words, 2.63 Cognitive Processes, 5.92 Allure, 11.84 Moralization, 0.00 Summary Variables Analytic, 58.24 Authentic 70.27

Apathy at 12

"The problem of a high pressure pump. I have much to do, and much to do. I have much to do, and much to do. I have much to do, and much to do. I have much to do, and much to do. I have much to do, and much to do. Why do you want to keep the a day. I've got a headache. Wish I had a soul. Wish I had a soul. You see much of a difference between you and me. When you're in the world. I miss no good. It's no good. It's no good. I've got a body. If I got any time to myself. I go out. I'm with you. I didn't like to get out. I had a heart. You're in the world. It didn't work. You're in the world. You can get out. You..."

I, me, my, 12.60 Positive Tone, 3.15 Negative Tone, 0.00 Social Words, 8.66 Cognitive Processes, 7.87 Allure, 17.32 Moralization, 0.00 Summary Variables Analytic, 17.02 Authentic 84.10

Apathy at alpha 12

"As for my daily life. I have a lot of time for a lot of people. I got a job. I have a boyfriend, got a husband. My parents care about me. They don't care about me. I got a house. I've seen the family really. I got much money. It really matters for me. I have to work. I got no friends. I got a phone. I really got no friends. I got a phone. I was really sick I got a phone. I have a phone. I really got a phone. You get much money I have a phone. You get much money. I know. A lot of people got real, that's what they got. You got much money. Nobody get much money. You got much money. I got much money, that's what I got."

I, me, my, 17.52 Positive Tone, 1.46 Negative Tone, 0.73 Social Words, 8.03 Cognitive Processes, 7.30 Allure, 17.52 Moralization, 0.00 Summary Variables Analytic, 55.91 Authentic 66.30

Apathy at alpha 12

"I'm here to get no one and I'm here to get no one." "My family even!" The family!" I am here to get no one. My family even! My personal life even. I go out to have any of the things I can take care of. I'm a victim of my own, so I sit here. We are in public. We've got a problem. I was born, so I've got a problem with that. We're in public. We've got a problem. I'm here to get no one. My family, even. I go out to have any of the things I can take care of. I'm a victim of my own, so I sit here."

I, me, my, 17.54 Positive Tone, 1.75 Negative Tone, 0.00 Social Words, 6.14 Cognitive Processes, 6.14 Allure, 20.18 Moralization, 0.00 Summary Variables Analytic, 1.63 Authentic 98.88

Apathy at alpha 15

"You get to keep coming. No This. That. Void of a belief. A very cold Much of a There. How. What Because. He Bored. If GOD. He I. B. He Fun. We. Neither We. Neither He B. B. We. Neither. Void. We. We. Neither Void. We. We. B. B. We. Neither. Void. We. We. B. B. We. None. Void. We."

I, me, my, 1.67 Positive Tone, 3.33 Negative Tone, 1.67 Social Words, 28.33 Cognitive Processes, 15.00 Allure, 6.67 Moralization, 0.00 Summary Variables Analytic, 1.12 Authentic 74.01

Apathy at 15

"I play games. I even go to watch. I am not really into playing with a brain myself, I really get no matter, so I am all that. I am that brain anyway. I am really like a person, I care about. What I am here for. I watch much, when I am in any. My brain. I want. I want to. I care. My brain. I get. I got around. My brain. I care. I care either. I play with the brain. I get. So. I am all that. I was, I am a. I got around. I got around. I am a brain. I got. I am. I am. I am a. I care. I am. I am. I care. I am."

I, me, my, 31.45 Positive Tone, 7.26 Negative Tone, 0.00 Social Words, 0.81 Cognitive Processes, 8.06 Allure, 5.65 Moralization, 0.00 Summary Variables Analytic, 0.47 Authentic 99.99

Peace at alpha 15

"The following is, in very long form, the official email you requested, from the person who is here at the North American Aerospace, I am talking about, among others, I am here at the I am at, "I am here, for the I am, a here, a a you, I am here a I am here a, is a I am here the, "the most an here. We are here for the I am, at, the I am and I am at here, we are here we are here, here the I am the I am, you are here, for the I am is, a I am here, indeed, I am here, I am here at, the I am and I am here, it is today. We are here, the I am, is not the I am, that is a here is, the most important. The I am, the I am here, the I am is here..."

I, me, my 15.29 Positive Tone 0.64 Negative Tone 0.00 Social Words 5.73 Cognitive Processes 1.91 Allure 15.29 Moralization 0.00 Summary Variables Analytic 5.55 Authentic 99.77

Peace at alpha 15

"Soliton. Here the earth is slowly sinking. It is a quiet night. It is a calm down, I am not asleep! I am talking to you. You are watching my words. It is a beautiful star. These are the words I'm talking about... The Earth here is now about to be here. This is the best I am going to have. What is this? I am asking you to be like me. Listen to me. You are reading this, I am talking to you. You are now listening to my words, I am talking to you. These are my words. I am speaking to you in a nice place, here's where it is. The place you will end. So listen to me. That's all I am saying, I am talking. Is this the way I am saying this? This is the way it is happening, a wonderful story. I am reading this, I am talking to you. I..."

I, me, my 12.42 Positive Tone 2.48 Negative Tone 0.00 Social Words 6.21 Cognitive Processes 0.62 Allure 8.07 Moralization 0.00 Summary Variables Analytic 5.29 Authentic 86.60

Peace at alpha 12

"Is anyone interested in a discussion?" she asked the others. "The discussion is for our benefit," I said, "and I don't ever recall agreeing with anybody during the meeting." "The only way you are going to become better," I added, "we are going to be making a lot of decisions about you over the coming months." "It is like that only I will be happier if we will never stop until we have completed this. I am not going to be content with this," I added. "It is like that only I am going to be content if we will never stop excepting your comfort," she said, "it will be like that only I have no choice except it will be like that only we will be content with all of this," I added, and in between the group...

**I, me, my 6.43 Positive Tone 3.57 Negative Tone 0.00 Social Words 10.00 Cognitive Processes 8.57 Allure 7.14
Moralization 0.00 Summary Variables Analytic 29.75 Authentic 72.59**

Peace at alpha 12

"I am a writer living in the present, working in a time where I am only going to get more excited about the future. I find myself in a weird spot when I think I have a moment. It is only a moment, and it has been too long. But where did I come from? Where am I going? These were the only words I had the thought before. Forget those fancy, "The place was pretty. You know, this feels good and this feels comfortable and that's what I have to do and I will be there for you." I was actually sitting in the chair. Where? Oh, in a chair. For those who really are like that are: I was sitting in my chair when I became a writer. I thought. What is that? No. But you know when you think a person is pretty? If they are, they might as well not be. That is just a really wonderful word..."

**I, me, my 9.82 Positive Tone 3.68 Negative Tone 0.00 Social Words 6.13 Cognitive Processes 12.88 Allure 8.59
Moralization 0.00 Summary Variables Analytic 4.70 Authentic 99.77**

Peace at alpha 10

Soybean, or Glycine max (Maysville), is a model plant in that it requires a consistent and accurate model for a long period of existence. When you own a garden, you don't have to look, and you don't mind paying an extra cost in return, if you choose to allow plants, as a guide, to thrive within your garden or the place of residence. For any kind of plant, the seeds you choose, for example, in the shape of the seed, are made up of the most expensive, the least expensive, the perfect food for the seed that is growing on a plant which is very cheap, in addition to the weather. So, if you have to look around your garden or you're looking to put a plant in your residence, then it has got to be you. A place which is very cheap, a very cheap, a very cheap.

**I, me, my 0.00 Positive Tone 1.32 Negative Tone 0.00 Social Words 7.24 Cognitive Processes 11.18 Allure 5.26
Moralization 0.00 Summary Variables Analytic 94.00 Authentic 52.18**

Peace at alpha 10

"Cupcake! Posted by Michael Sargen, Wednesday September 9, 2012 I thought that I was going to like this recipe for the "cupcake"! It came out well-chilled, so I went to the kitchen at my family's vacation get-together this week and decided to enjoy one of the best cupcakes I have ever tasted. However, I didn't get that great of a result, so I did some serious blogging, and here is what I learned: It was a perfect cupcake. Perfect cupcake, as they say. I really enjoyed my cupcakes (I really really really mean that) and I can't wait to try these other recipes that were mentioned in my last post. Here's what I learned: 1. I made these cupcakes in my favorite 9-by-9-fluted cake pan. 2. I baked a whole batch of the cake as..."

**I, me, my 11.97 Positive Tone 4.23 Negative Tone 0.00 Social Words 0.70 Cognitive Processes 11.97 Allure 9.15
Moralization 0.00 Summary Variables Analytic 38.48 Authentic 88.19**

Peace at alpha 7

"in the early 1960s, the city of Toronto experienced what it called its "golden era," when the city was growing economically, having the highest gross domestic product in Canada. The city and other towns with gold in their city-planning history were being formed by population increases; by 1970, Toronto's population had reached 5 million. Toronto had also been known as an ideal place to build because it was the oldest and most developed urban area in North America. When the city was first founded in 1804, Toronto was known as York, where Toronto today. At this stage, the city was a great place to locate in the heart of the vast forest that was to become Toronto. A small village at the time of its birth, the city was created from Toronto, and that was the point at which the city was built. In fact, by the time the city was founded, the location where the City would later be built was a long walk from a..."

**I, me, my 0.00 Positive Tone 0.60 Negative Tone 0.60 Social Words 1.79 Cognitive Processes 5.95 Allure 4.17
Moralization 0.00 Summary Variables Analytic 92.45 Authentic 87.30**

Peace at alpha 7

“The world is a dangerous place for Christians - but also has many exciting things to share! I've said many times that I'd like to offer a place for Christians to find their place here, but I don't have a place for myself to find my place, or anyone else's place. This is because I'm the church you have to follow! Here is what I said in a Facebook post: My pastor is the same pastor I mentioned in the other thread about being Christian, a guy who calls himself Jesus Christ. In the other thread, we were referring to ourselves as Christians, but he was a nice, lovely person who I didn't feel like having a place for him. But this is his place, his church. You know that's not true when you say my church is just a place for Christians to go. You said that before. You know what I know? That's what I believe about it. But my pastor is Christian.”

**I, me, my 9.15 Positive Tone 3.05 Negative Tone 1.83 Social Words 12.80 Cognitive Processes 15.24 Allure 12.80
Moralization 0.00 Summary Variables Analytic 13.52 Authentic 97.21**